

Data Structures and Algorithms

Conf. dr. ing. Guillaume Ducoffe

`guillaume.ducoffe@fmi.unibuc.ro`

Operations on/with Trees

Queries on Rooted Trees:

- Least-common ancestors
- Distances

Tree Decompositions Methods

- Heavy-Path (HP)
- Centroid

Tree-Based Data Structures:

- Cartesian Trees

Least common ancestor (LCA) queries

Every two nodes x, y have a common ancestor (e.g., the root)

Definition


$LCA(x, y) =$ **deepest** common ancestor (furthest from the root)

Examples:

$$LCA(15, 7) = 6$$

$$LCA(15, 12) = 4$$

$$LCA(15, 9) = 8$$



Pre-processing time? Query time?

Naive resolution

While nodes x and y are distinct, replace the deepest node by its parent.

```
node *lca(node *x, node *y) {  
    while(x != y){  
        if(x->level >= y->level)  
            x = x->father;  
        else y = y->father  
    }  
    return x; //==y  
}
```

Pre-processing: $\mathcal{O}(n)$ – Computation of the levels

Query Complexity: Linear in the **height** of the tree

Extremal cases

- Stars have height = 1



- But Paths (\sim Lists) have height = $n - 1$



- Bounded-degree trees (e.g., Binary) have height in $\Omega(\log n)$




Heavy-path decomposition

- For each non-leaf node, select a child with maximum number of descendants in its rooted subtree
 - variant: select the unique child y of x s.t. $y.order > x.order/2$, if any.

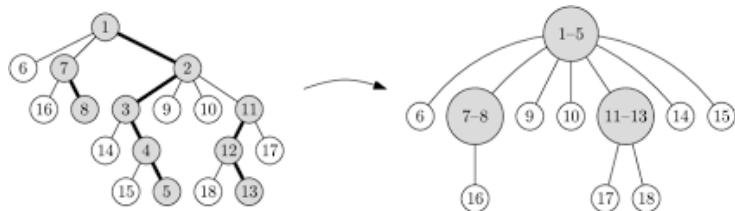
⇒ Partition of the nodes in so-called “heavy” paths.

HP-tree: rooted subtree obtained by contracting each HP in one node.



Encoding

- We can associate to each node a number (ID of its HP).
- We can store in an array the highest/deepest node of each HP.



- We may also associate, to each node, a pointer to its child in the same HP (possibly, null).

Implementation

//k denotes the ID of the current HP
//first and last store ends of each HP
//hp denotes HP, next_hp points to next child in HP

```
void compute_hp(node *n, int& k, vector<node*>& first, vector<node*>& last){
    n->hp = k; last[k] = n;
    if(n->child != nullptr){
        node *c = n->child;
        for(node *p = c->next; p != nullptr; p = p->next)
            if(p->order > c->order) c = p;
        n->next_hp = c; compute_hp(c,k,first,last);
        for(node *p = n->child; p != nullptr; p = p->next)
            if(p != c) { //new HP
                first.push_back(p); last.push_back(p);
                compute_hp(p,++k,first,last);
            }
    }
}
```

Complexity: $\mathcal{O}(n)$

LCA queries with HP

Logarithmic-time version

- Compute levels in the HP-tree.

```
void levels_hp(node *n, vector<int>& hp_lvl) {
    if(first[n->hp] == n) {
        if(n->father == nullptr) hp_lvl[n->hp] = 0;
        else hp_lvl[n->hp] = hp_lvl[n->father->hp] + 1;
    }
    for(node *c = n->child; c != nullptr; c = c->next) { levels_hp(c, hp_lvl); }
}
```

- Simulation of the naive algorithm on the HP-tree.

//for ease of writing, we assume first and last to be global arrays

```
node *lca_hp(node *x, node *y) {
    while(x->hp != y->hp){
        if(x->level >= y->level)
            x = ( (x == first[x->hp]) ? x->father : first[x->hp] );
        else y = ( (y == first[y->hp]) ? y->father : first[y->hp] );
    }
    return ( (x->level >= y->level) ? x : y );
}
```

Property: The HP-tree has height in $\mathcal{O}(\log n)$.


Improvements: Binary Lifting

Binary tree of height in $\mathcal{O}(\log n)$.

1) Label the root with 1. If a node is labelled with i , then label its left (resp. right) child by $2i$ (resp., $2i + 1$).

→ Keep the pairs (label,node) in a hash table

2) A node at level i has a $(i + 1)$ -bit label. The j^{th} closest ancestor of a node x has label $x.label \gg j$ (division by 2^j)



Remark: if the LCA of two nodes x, y is at level i , then the $i + 1$ most significant bits of their respective labels are identical.

Binary Lifting: Application to LCA

Input: nodes x, y

s.t. $x.level \leq y.level$.

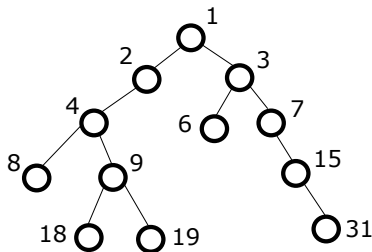
1) $a = x.label$ and $b = y.label$

2) If $i = y.level - x.level$ is > 0 , then
 $b = b \gg i$ (division by 2^i).

3) Let $c = a \oplus b$ (bitwise XOR)

4) Let $j = \lfloor \log c \rfloor + 1$ (most significant bit)

5) $LCA(x, y)$ is the node with label $a \gg j$.




Complexity: $\mathcal{O}(1)$ assuming bitwise operators (but $\mathcal{O}(\log n)$ otherwise).

Intermezzo: A problem of **intervals**

Consider a pre-order of the nodes.

Observation: a rooted subtree = an interval



$LCA(x,y)$, $x < y \implies$ **shortest** interval that contains $[x,y]$

DFS computation

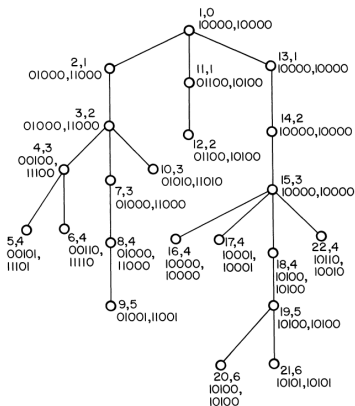
```
void compute_intervals(node *n, int& num){  
    n->start = num++;  
    for(node *c = n->child; c != nullptr; c = c->next)  
        compute_intervals(c,num);  
    n->end = num;  
}
```

Complexity: $\mathcal{O}(n)$

Remark: Post-ordering = nodes ordered by increasing end value (break ties by decreasing height...)

Binary Lifting + HP-tree

Pre-processing



1) Compute a preorder + preorder-intervals [start,end] for each node

2) For each node v , let its **inlabel** denote the node in its subtree whose preorder number has maximum number of rightmost 0s in its binary representation.

3) **HP decomposition using inlabels**
(same HP = same inlabel)

Pre-processing time: $\mathcal{O}(n)$ + computation of inlabels

Computation of inlabels

- We compute, for each node v , the number of rightmost 0s in $v.preorder$.

→ For that, we simulate the increment of a binary counter with a stack.

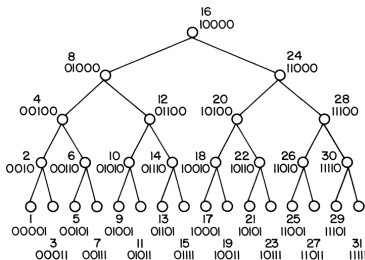
- We consider the nodes in preorder
- At the time we consider the i^{th} node, a stack S memorizes the position of all 1's in decreasing order (from right to left) in the binary representation of $i - 1$.

Incrementation: pop until we find two nonconsecutive 1's or S becomes empty. **Amortized complexity** $\mathcal{O}(1)$

- Number of rightmost zeros = position of the rightmost zero (access to top element)
- Inlabels by dynamic programming on the tree. $\implies \mathcal{O}(n)$

Binary Lifting + HP-tree

Properties



1) Identify the inlabels with the **inorder** numbers of some nodes in the smallest **complete binary tree** B with $> n$ nodes.

2) **Descendance-preservation property**: if x is a descendant of y , then $x.inlabel$ is a descendant of $y.label$ (the converse is false in general)

3) In a complete binary rooted tree of order $2^{h+1} - 1$, for a node x at level i , we can compute its j^{th} closest ancestor by suppressing the $i, i-1, \dots, i-j+1$ most significant bits and adding 0s to the right.

Consequences: if $z = LCA(x, y)$, then $z.inlabel$ is an ancestor of $w = LCA(x.inlabel, y.inlabel)$ in the complete binary rooted tree B .

Furthermore, we can compute w using binary lifting.

Binary Lifting + HP-tree

Query (1/2)

Input: nodes x, y with $x.inlabel \neq y.inlabel$.

1) Compute $LCA(x.inlabel, y.inlabel)$ in B

- Compute the difference of level between $x.inlabel$ and $LCA(x.inlabel, y.inlabel)$
(from the most significant bit of $x.inlabel \text{ XOR } y.inlabel$)
- Binary lifting: Compute $a = (x.inlabel \ll j) \gg j$ (most significant bits), $b = (x.inlabel - a)$ then $c = b - ((b \ll k) \gg k)$ (suppression of the bit subsequence). Output $a + c$.

Remark: requires level of $x.inlabel$ in B , that can be deduced from the number of rightmost 0s (already computed)

In what follows, let $w = LCA(x.inlabel, y.inlabel)$

The **ascendant** numbers

Reminder: for every node r , we know the level of $r.inlabel$ in B

- We define $r.ascendant$ so that its j^{th} bit is set to 1 if and only if it has an ancestor s in T whose $inlabel$ is at level j in B .

We can compute the ascendant numbers by dynamic programming on the tree T .

Sketch: Let p denote the parent of r . We obtain $r.ascendant$ from $t.ascendant$ by setting one new bit to 1 (namely, the level of $r.inlabel$ in B). This can be done using bitwise operators:

$$r.ascendant = t.ascendant + (1 \gg i)$$

Additional pre-processing time: $\mathcal{O}(n)$

Binary Lifting + HP-tree

Query (2/2)

Input: nodes x, y with $x.inlabel \neq y.inlabel$.

2) Let $z = LCA(x, y)$ (we do not know z at this point). In order to compute $z.inlabel$ from w , we compute the level of $z.inlabel$ in B .

→ XOR on $x.ascendant, y.ascendant$

3) To compute z , we need to find the closest ancestors of x, y in the heavy-path corresponding to $z.inlabel$. Let us detail for x . We assume $x.inlabel \neq z.inlabel$. Let P_x, P_z denote the two HPs of x, z . We just need to find the neighbour of P_z on the unique $P_z P_x$ -path in the HP-tree.

→ The inlabel of this HP can be retrieved from $x.ascendant$! (again using bitwise operators).

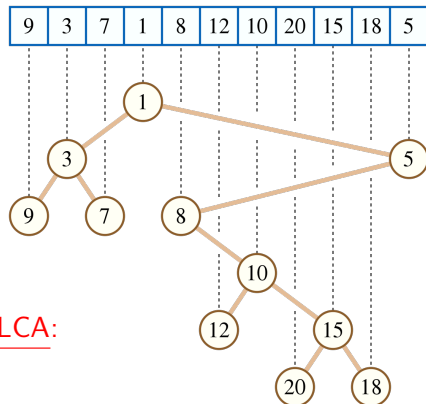
Query time: $\mathcal{O}(1)$ using bitwise operators, and $\mathcal{O}(\log n)$ otherwise.

An application to Range queries: Cartesian trees

Input: an n -size vector $v[]$

Output: an n -node binary rooted tree T whose nodes are the elements of $v[]$, such that:

- The root is a min/max element
- the left/right subtrees are Cartesian trees for the left/right subvectors to the root.



Reduction of min/max range queries to LCA:

The min/max elements between indices i and j is the element in position $LCA(i,j)$.

Construction

- It suffices to compute the vector `father` encoding the parent of each node.
- We scan the vector from left to right and maintain potential candidates for parent nodes in a stack. We repeatedly pop out smaller/bigger elements. \implies **left neighbour**

Interpretation: if the father of a i is at an earlier position $j < i$, then it must be its left neighbour.

- Compute the **right neighbours** in the same way by scanning from right to left.
- Father of a node: the largest/least value amongst its left and right neighbours.

Implementation

Min. version

```
vector<int> compute_cartesian_tree(const vector<int>& v) {  
    vector<int> left(v.size());  
    stack<int> candidates;  
    for(int i = 0; i < v.size(); i++) {  
        while(!candidates.empty() && v[candidates.top()] > v[i])  
            candidates.pop();  
        left[i] = (candidates.empty()) ? -1 : candidates.top();  
        candidates.push(i);  
    }  
    /* right neighbours */  
    vector<int> father(v.size());  
    for(int i = 0; i < v.size(); i++) {  
        father[i] = (left[i] <= right[i]) ? right[i] : left[i];  
    }  
    return father;  
}
```

Complexity: $\mathcal{O}(n)$ (Potential: size of the stack)

Beyond LCA: Distance queries

- Distance between two nodes x and y : number of edges on the (unique) path in the tree between x and y
- Application: Routing in Tree Networks (or in general networks using a spanning trees)
- All distances can be computed in $\mathcal{O}(n^2)$ time by varying the root and applying BFS.

Can we do better?

Reduction to LCA

Input: nodes x, y .

1) Let $z = LCA(x, y)$.

2) Node z must be on the xy -path

$$\implies d(x, y) = d(x, z) + d(z, y).$$

3) Since z is an ancestor of x (resp., y), we have

$$d(x, z) = x.height - z.height \text{ (resp., } d(y, z) = y.height - z.height).$$

Pre-processing: $\mathcal{O}(n)$ (compute the heights + LCA pre-processing)

Query time: $\mathcal{O}(1)$

Centroids in trees

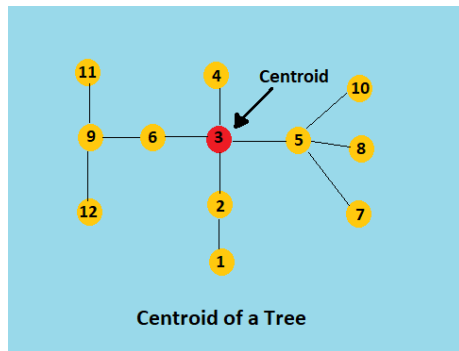
Definition

A centroid in an n -node tree T is a node c such that every subtree of $T \setminus \{c\}$ has order $\leq n/2$.

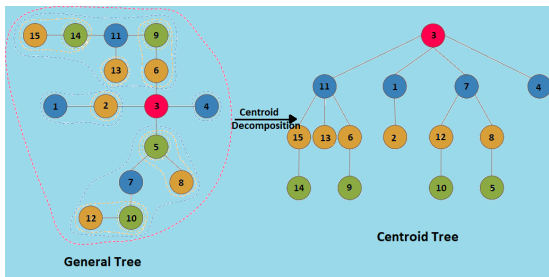
A centroid always exists.

Proof:

- 1) There are $n - 1$ edges (all nodes but the root have a parent)
- 2) For each node n , choose a heaviest subtree of $T \setminus \{n\}$ and orient the edge from n to this subtree.
- 3) One edge is oriented in both directions!



Centroid decomposition



- 1) Compute a centroid c
- 2) Compute a centroid decomposition for each subtree of $T \setminus \{c\}$
(output = rooted tree)
- 3) Merge all decompositions in one rooted tree with c as root.

The result is a rooted tree T' with same node-set as the original tree T .

Property: the centroid decomposition outputs a tree of height in $\mathcal{O}(\log n)$

Application to LCA/Distance queries: store for each node its path-to-root in the centroid decomposition + distances in T .

Computation

- A centroid can be computed in $\mathcal{O}(n)$ time.
- Pre-compute the order of each rooted subtree.
- **Local search.** Start from any node. If each subtree of $T \setminus \{c\}$ has order $\leq n/2$ then output c . Otherwise, go to a neighbour in a heaviest subtree.

Remark: for each child c' of c we know the order of its rooted subtree. If c' is the parent of c , then the subtree of $T \setminus \{c\}$ that contains c' has order $n - c - > \text{order}$.

- There are $\mathcal{O}(\log n)$ recursive steps
- \implies Centroid decomposition in $\mathcal{O}(n \log n)$ time.

Questions

