

Data Structures and Algorithms

Conf. dr. ing. Guillaume Ducoffe

`guillaume.ducoffe@fmi.unibuc.ro`

Searching an element

One of the most basic problems in Computer Science.

- In an unordered structure: $\mathcal{O}(n)$ time.
- In an ordered vector: $\mathcal{O}(\log n)$ time.
- In an ordered skip list: **expected** $\mathcal{O}(\log n)$ time.
- In a hash table: **expected** $\mathcal{O}(1)$ time.

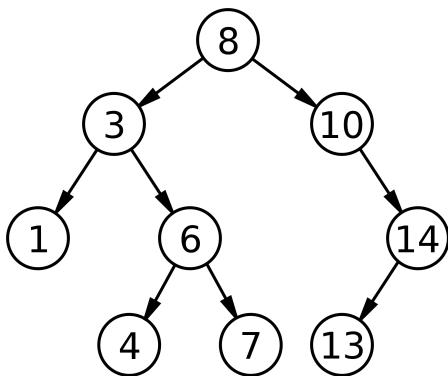
(binary search)

Goal: comparable performances with a dynamic + deterministic data structure.

Binary Research Tree

Each node stores a value i .

- Nodes in the left subtree store values $< i$.
- Nodes in the right subtree store values $> i$.



Remark: we may assume that all values stored are pairwise different (just add a counter for the number of occurrences of each element)

Implementation

Just a binary tree...

```
struct node {  
    int value;  
    node *father, *left, *right; //sometimes father is omitted  
};  
typedef node *BinaryResearchTree;
```

- Four standard operations:
 - Emptiness test
 - Search for an element
 - Insertion of an element
 - Deletion of an element

Operations (1/3)

Naive Implementation

//Complexity: $\mathcal{O}(1)$

```
bool empty(const BinaryResearchTree& T) { return (T==nullptr); }
```

//Can be modified to output any desired information: node pointer, height, etc.

//Complexity: $\mathcal{O}(\text{height})$

```
bool search(const BinaryResearchTree& T, int e) {  
    if(empty(T))  
        return 0;  
    else if(T->value == e)  
        return 1;  
    else if(T->value < e)  
        return search(T->right,e);  
    else  
        return search(T->left,e);  
}
```

Operations (2/3)

Naive Implementation

//Complexity: $\mathcal{O}(\text{height})$

```
void add(BinaryResearchTree& T, int e) {  
    if(empty(T)) {  
        T = new node;  
        T->father = T->left = T->right = nullptr;  
        T->value = e;  
    }else if(e < T->value) {  
        add(T->left,e); if(!empty(T->left)) T->left->father = T;  
    }else if(e > T->value) {  
        add(T->right,e); if(!empty(T->right)) T->right->father = T;  
    }//nothing to be done if e == value  
}
```

Remark: if multiple occurrences allowed, put all equal elements to left/right.

Operations (3/3)

Naive Implementation

//Complexity: $\mathcal{O}(\text{height})$

```
void remove(BinaryResearchTree& T, int e) {
    if(!empty(T)) {
        if(e < T->value) remove(T->left,e);
        else if(e > T->value) remove(T->right,e);
        else { //e == value
            if(empty(T->left) && empty(T->right)) { //emptied tree
                BinaryResearchTree tmp(T);
                T = nullptr; delete tmp;
            } else if(!empty(T->left)) {
                T->value = maximum(T->left);
                remove(T->left, T->value);
            } else {
                T->value = minimum(T->right);
                remove(T->right, T->value);
            }
        }
    }
}
```

Remark: min/max is called at most once. Removing the min/max of a subtree can be done in $\mathcal{O}(1)$ because it is a leaf node.

Min/Max element

//Complexity: $\mathcal{O}(\text{height})$

```
int minimum(const BinaryResearchTree& T) {  
    if(empty(T->left)) return T->value;  
    else return minimum(T->left);  
}
```

//Complexity: $\mathcal{O}(\text{height})$

```
int maximum(const BinaryResearchTree& T) {  
    if(empty(T->right)) return T->value;  
    else return maximum(T->right);  
}
```


Generalization: Order statistics

Definition (k^{th} order statistic)

The k^{th} smallest element.

Special cases:

- $k = 1$ (Minimum) and $k = n$ (Maximum)
- $k = n/2$ (if n even): **Median**

If n odd then we have Lower/Upper Median, and the Median equals their average

- $k = n/4, n/2, 3n/4$ (Quartiles), etc.

Computation

Every node stores in some integer variable the **size of its left subtree**.

→ To be updated after each insertion/deletion.


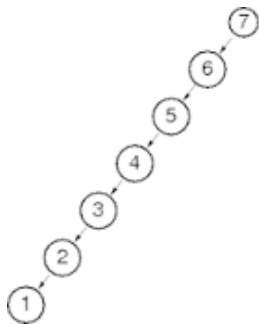
//Complexity: $O(\text{height})$

```
int stat(const BinaryResearchTree& T, int k){  
    if(T->leftSize==k-1) return T->value;  
    else if(T->leftSize >= k) return stat(T->left,k);  
    else return stat(T->right,k-1-T->leftSize);  
}
```

Remark: slight changes needed if multiple occurrences allowed...

The height of a Binary Research Tree

- In the worst-case: $\mathcal{O}(n)$.
- In the best-case: $\mathcal{O}(\log n)$.



Objective: Keep the height to $\mathcal{O}(\log n)$ (**Balanced Tree**).

Balanced Tree: Offline Construction

Key observation: A Binary Tree is balanced if for every node with p descendants:

- $\leq c \cdot p$ nodes are in its left subtree;
- $\leq c \cdot p$ nodes are in its right subtree;

for some constant $c \in [1/2; 1)$.


- Natural choice: $c = 1/2 \implies$ Each subtree has for root its **median** value!

Algorithmic problem: fast computation of the median (or of the cn^{th} -order statistic, for some $c \in [1/2; 1)$)?

Statistic search: Divide & Conquer strategy

Quickselect

- 1) Select an arbitrary element p (sometimes called **pivot**) and partition in two sub-vectors: one for smaller elements, and one for larger elements.
- 2) If element p is not the k^{th} order statistic, then recurse on one of the two sub-vectors.



Implementation

```
//Returns the final position of the pivot element
int partition(vector<int>& v, int lb, int ub) {
    int e = v[lb]; //to be discussed...
    int p = ub; //position of the pivot element
    for(int i = ub; i > lb; i--)
        if(v[i] >= v[lb]) {
            swap(v,i,p--); //new larger element found
        }
    swap(v,lb,p);
    return p;
}

int quickselect(vector<int>& v, int k, int lb, int ub) {
    int p = partition(v,lb,ub);
    if(p==k-1) return v[p];
    else if(p >= k) return quickselect(v,k,lb,p-1);
    else return quickselect(v,k,p+1,ub);
}
```

Complexity

- Fix the n elements of the vector and consider a **random permutation**.

→ Every element is put in position 0 with same probability $\frac{1}{n}$. In particular, the left sub-vector has length k with probability $\frac{1}{n}$.

→ The **Average complexity** satisfies the following inequation:

$$T(n) \leq n - 1 + \sum_{k=1}^{n-1} \frac{1}{n} \times \max\{T(k), T(n-1-k)\}$$

By induction: $T(n) \leq 4 \cdot n$.

Alternative strategy: choose a **random pivot**.

Expected complexity in $\mathcal{O}(n)$.

Median of medians

Key observation: To ensure Linear-time **Worst-case** complexity, it would suffice to choose as pivot a cn^{th} order statistic, for some $0 < c < 1$.

$$\begin{aligned} T(n) &\leq n + T(cn) \leq n + cn + T(c^2n) \leq \dots \leq n + cn + c^2n + \dots + c^i n + \dots \\ &= n \times \sum_i c^i = \mathcal{O}(n) \end{aligned}$$

- We subdivide the vector in sub-vectors of size 5 and we keep the median of each of them.

[1, 4, 5, 22, 9, 13, 67, 91, 0, 15, 33, 50, 16, 12, 87, 19, 14] \implies [5, 67, 50, 14]

The new vector has size $\leq n/5$ (*Sufficiently small for recursive calls*). Its median is the cn^{th} order statistic of the original vector, for some $c \in [3/10; 7/10]$.

Sketch Implementation

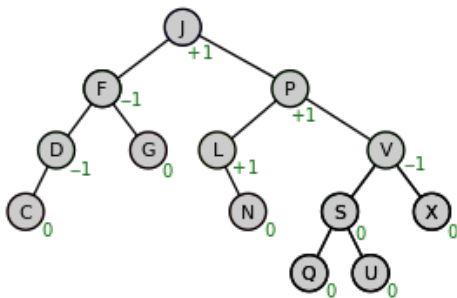
```
int quickselect(vector<int>& v, int k, int lb, int ub) {  
    vector<int> med;  
    for(int i = lb; i <= ub; i+=5) {  
        med.push_back(quickselect(v,i+2,i,i+4));  
        //border effects for last subvector...  
    }  
    int p = quickselect(med,med.size()/2,0,med.size()-1);  
    //Find p in the vector, then use it as pivot  
    ...  
}
```

Complexity: $T(n) \leq n - 1 + T(n/5) + T(7n/10) = \mathcal{O}(n)$.

Online construction: AVL trees

Adelson-Velsky and Landis (1962).

- Every node keeps track of the order of its **height**
- After each operation, we modify these subtrees so that $|left.height - right.height| \leq 1$.



Every AVL of height h contains $\geq F(h)$ nodes (by induction). \implies

Balanced

AVL trees

Implementation

```
struct node {  
  
    int value;  
  
    node *father, *left, *right;  
  
    int height;  
  
};  
  
typedef node *AVL;
```

In practice, we may only store `left.height - right.height` (only 2 bits needed). However, various complications would arise after each insertion/deletion.

Actualizing the height value

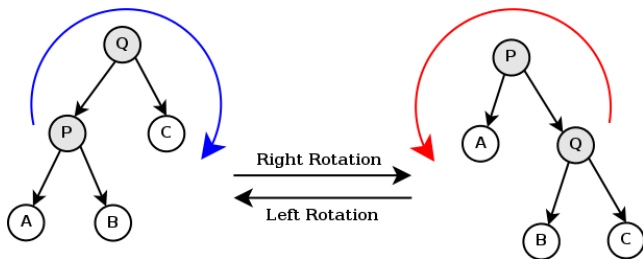
Just to have it somewhere for future reference ...

```
void update(AVL& T) {  
    if(!empty(T)) {  
        if(empty(T->left) && empty(T->right))  
            T->height = 0;  
        else if(empty(T->left))  
            T->height = 1 + T->right->height;  
        else if(empty(T->right))  
            T->height = 1 + T->left->height;  
        else if(T->left->height >= T->right->height)  
            T->height = 1 + T->left->height;  
        else T->height = 1 + T->right->height;  
    }  
}
```

Tree rotation

The left/right child becomes the new root.

Left/right grandchildren need to be redistributed. \implies We need to use the pointer to parent



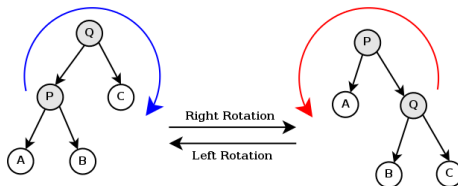
Tree rotation

Motivation

Consider the following scenario:

- subtree(P) and subtree(C) are balanced
- $P.height = C.height + 2$ (subtree(Q) is not balanced)
- $A.height = P.height - 1$ (highest subtree) = **C.height + 1**

After a right rotation, the tree becomes balanced!



Tree rotation

Implementation

```
void rotateRight(AVL& T) { //Assumption: left is nonempty
    AVL newRoot = T->left; newRoot->father = T->father;
    if(!empty(T->father)){
        if(T->father->left == T) T->father->left = newRoot;
        else T->father->right = newRoot; }
    T->left = T->left->right; if(!empty(T->left)) T->left->father = T;
    newRoot->right = T; T->father = newRoot;
    update(T); update(newRoot);
    T = newRoot; }

void rotateLeft(AVL& T) { //Assumption: right is nonempty
    AVL newRoot = T->right; newRoot->father = T->father;
    if(!empty(T->father)){
        if(T->father->left == T) T->father->left = newRoot;
        else T->father->right = newRoot;}
    T->right = T->right->left; if(!empty(T->right)) T->right->father = T;
    newRoot->left = T; T->father = newRoot;
    update(T); update(newRoot);
    T = newRoot; }
```

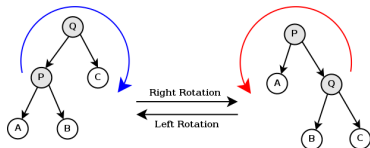
Adding a value

Suppose we add a new element **to the left**

⇒ We may have $\text{left.height} = \text{right.height} + 2$ (Not balanced)

- Case 1: the left-left subtree is higher than the left-right subtree

→ we only need one rotation




Adding a value

Suppose we add a new element **to the left**


⇒ We may have $\text{left.height} = \text{right.height} + 2$ (Not balanced)

- Case 1: the left-left subtree is higher than the left-right subtree

→ we only need one rotation



- Case 2: Left rotation (make left-left subtree heavier + AVL) then Right rotation (as for Case 1)



Adding a value

Implementation

```
void add(AVL& T, int e) {
    if(empty(T)) { /*already discussed*/ }
    else if(e == T->value) return;
    else if(e < T->value) {
        add(T->left,e);
        if(!empty(T->left)) {
            T->left->father = T;
            int r = (empty(T->right)) ? 0 : T->right->height;
            if(T->left->height == r+2) {
                if(!empty(T->left->left) && T->left->left->height == r+1) //Case 1
                    rotateRight(T);
                else { //Case 2
                    rotateLeft(T->left); rotateRight(T);
                }
            }
        }
    } else { /*add to the right*/ }
    update(T);
}
```

Removing a value

- Case we remove an element **to the right**
 \implies We may have $\text{left.height} = \text{right.height} + 2$ (Not balanced)

Proceed as before. . .

- Case we remove an element **to the left**: Same as above. . .

Removing a value

- Case we remove an element **to the right**
 \implies We may have $\text{left.height} = \text{right.height} + 2$ (Not balanced)

Proceed as before. . .

- Case we remove an element **to the left**: Same as above. . .
- Case we remove the root. The **new root** can be:
 - either the maximum to the left
 - or the minimum to the right

\implies We choose the element in the highest subtree.

Removing a value

```
void remove(AVL& T, int e) {
    if(!empty(T)) {
        if(e < T->value) {
            remove(T->left,e);
            //Do as for insertion
        } else if(e > T->value) {
            remove(T->right,e);
            //Do as for insertion
        } else { //e == value
            if(T->height == 0) { /* emptied tree: proceed as before*/ }
            else if(!empty(T->left) && T->left->height == T->height - 1) {
                T->value = maximum(T->left); remove(T->left,T->value);
            } else {
                T->value = minimum(T->right); remove(T->right,T->value);
            }
        }
        update(T);
    }
}
```


Beyond Binary Research Trees

- Research Trees of larger Arity
 - B-trees
 - 2 – 3 – 4 trees
- Research Trees for multi-dimensional points
 - Interval trees
 - k -range trees

Storing more values per node: B-trees

Standard Data Structure for indexing in Data Bases and File Systems.

- Two parameters L, U (in general, $U = 2L$):
 - Each internal node contains $\geq L - 1$ values;
 - Each node contains $\leq U - 1$ values.
- If the values stored in an internal node are a_1, \dots, a_k , $L - 1 \leq k \leq U - 1$ then there are exactly $k + 1$ subtrees:
 - Nodes with values $\leq a_1$
 - Nodes with values $> a_{i-1}$ and $\leq a_i$, $\forall 2 \leq i \leq k$
 - Nodes with values $> a_k$



Implementation

```
struct Bnode {  
    vector<int> values;  
    vector<Bnode*> children;  
    Bnode *father;  
};  
typedef Bnode *Btree;
```

Remark 1: If values is nonempty, then `children.size() == values.size() + 1`

Remark 2: `values.size()` lies between $L - 1$ and $U - 1$ (varying node sizes)

Remark 3: values is sorted so as to find efficiently (using binary search) the child node where to continue the search.

Path-balanced property

A B-tree must preserve the following invariant:

Definition

All leaves must stay at the same level (=distance to the root).

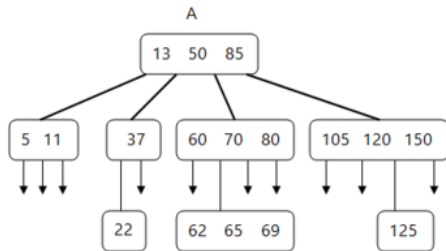
\implies The number of nodes grows by a factor $\geq L - 1$ at each step, and therefore there are at most $\mathcal{O}(\log n / \log L)$ levels.

Remark: path-balanced cannot be checked locally. We need operations that always preserve this property (*i.e.*, we cannot correct the tree if it becomes unbalanced after each operation).

2 – 3 – 4 trees

- Special case of B-trees for $L = 2, U = 4$.
- Can be implemented with Binary Research Trees!

Figure 1: Multiway Search Trees




Insertion of a value (1/3)

- Search for the element in the current tree
(we may assume that we did not find it...)

Example for $e = 67$

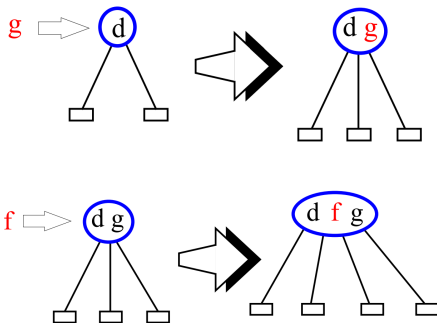
Figure 1: Multiway Search Trees



- Consider the last node on the search path (it is a leaf).

Insertion of a value (2/3)

- **Try to** insert the new element in the leaf node

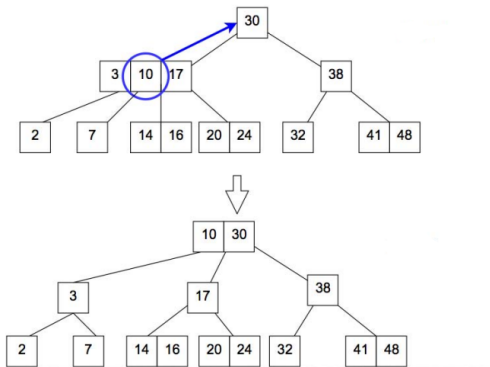


- Only possible if the node is not already full (4-node)

Insertion of a value (3/3)

Solution: Split all the 4-nodes on the search path!

The middle element is now stored in the parent node




Remark: preserves the path-balanced property...



Deletion of a value

- 1) Locate the value to be removed in some node N.
- 2) Ensure **recursively** that N contains at least two values.



- If the root contains one value, merge it with its two children.



- If the next node has one value *but* a closest sibling with > 1 values then make a transfer.



- Else, merge the node with a sibling and transfer one value from the parent.

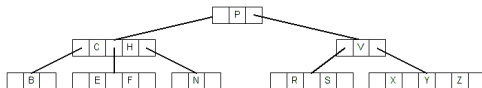
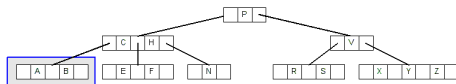


Deletion of a value

Case of a Leaf Node

Our pre-processing ensures that N contains at least one more element than the one to be removed.

⇒ Just remove the element



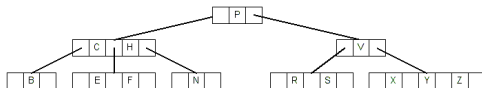
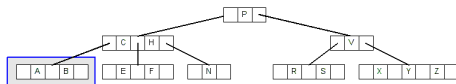
- What was the problem if there were exactly one value?

Deletion of a value

Case of a Leaf Node

Our pre-processing ensures that N contains at least one more element than the one to be removed.

⇒ Just remove the element



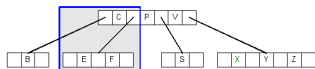
- What was the problem if there were exactly one value?

→ path-balanced property

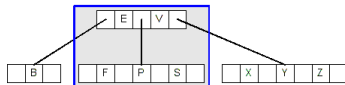

Deletion of a value

Case of an Internal Node

- Subcase # 1: One child has > 1 values
 \implies Replace the deleted value with its predecessor/successor.



- Subcase # 2: Both children have 1 value
 \implies Merge the two children with the value to be deleted.




Proceed **recursively**!

Encoding: Red-Black Trees

A 2 – 3 – 4-tree can be encoded as a **binary research tree**, with one colour (red/black) being assigned to each node.

Encoding of a Bnode:

- The middle value of a node is coloured black
- The (at most) two other values are coloured red.

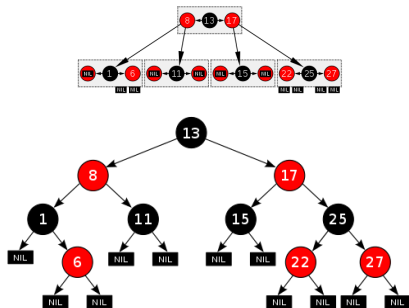


Application: simulate all operations on a 2-3-4-tree by operations on Binary Research Trees (mostly, **left/right rotations**).

Properties of Red-Black Trees

- Every node is either red or black
- The root is black
- Any child of a red node is black.
- There is the same number of black nodes on any root-leaf path

"all leaves of a 2 – 3 – 4 tree are at the same level"

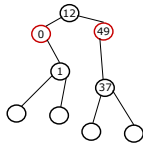
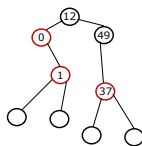


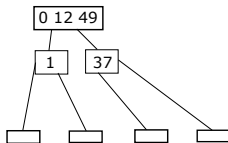
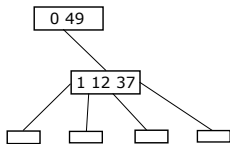


Implementation

```
struct RedBlackNode {  
    int value;  
    RedBlackNode *father, *left, *right;  
    bool color; //false for black  
};  
  
typedef RedBlackNode *RedBlackTree;
```

Splitting a 4-node

- Recolour black the two children nodes
- + At most two rotations
- + At most $\mathcal{O}(1)$ recolouring.



Splitting a 4-node

```
void checkAndSplit(RedBlackTree& T) {  
    if( (!empty(T->left)&& T->left->color)  
        && (!empty(T->right)&& T->right->color) ) {  
        //full node  
  
        T->left->color = T->right->color = 0; //black  
  
        if(!empty(T->father)) { //not the root  
            if(!T->father->color) { //black node  
                T->color = 1;  
            } else {  
                if(T->father->left == T) {  
                    rotateRight(T->father);  
                    T->right->color = 1;  
  
                    //new rotation  
                    if(T->father->left == T) {  
                        rotateRight(T->father); T->right->color=1;  
                    } else {  
                        rotateLeft(T->father); T->left->color=1;  
                    }  
                } else { /*right child*/ }  
            }  
        }  
    }  
}
```

More dimensions: **Interval tree**

Def.: an interval = an ordered pair of two values.

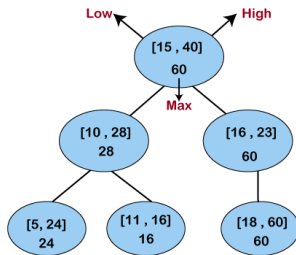
An **interval tree** maintains a dynamic collection of intervals, while supporting the following operations:

- Insertion/Deletion/Membership queries.
in $\mathcal{O}(\log n)$ time.
- Output an interval that contains a given point x .
in $\mathcal{O}(\log n)$ time.
- Output an interval that intersects a given interval $[a; b]$
in $\mathcal{O}(\log n)$ time.

Implementation: Augmented Tree

- We put all intervals in some **balanced** binary search tree (e.g., AVL)

→ Use lexicographic ordering



- Insertion/Deletion/Membership queries automatically in $\mathcal{O}(\log n)$ time.
- Additional storage of the largest upper values amongst all intervals in the subtree of a node.

Remark: other implementations exist with comparable performances.

Intersection with a point

Input: a point x

Output: any interval containing x .

- If $x \geq \text{root} \rightarrow \text{low}$ and $x \leq \text{root} \rightarrow \text{high}$ then output root.
- Else if $x < \text{root} \rightarrow \text{low}$, then we recurse on $\text{root} \rightarrow \text{left}$.
- Else, $x > \text{root} \rightarrow \text{high}$.
 - If $\text{!empty}(\text{root} \rightarrow \text{left})$ and $x \leq \text{root} \rightarrow \text{left} \rightarrow \text{max}$, then there is an interval on the left containing x . We recurse on $\text{root} \rightarrow \text{left}$.
 - Else, we recurse on $\text{root} \rightarrow \text{right}$.

Complexity: $\mathcal{O}(\log n)$.

Intersection with an interval

Input: an interval $[x; y]$

Output: any interval intersecting $[x; y]$.

- If $x \leq \text{root} \rightarrow \text{high}$ and $\text{root} \rightarrow \text{low} \leq y$ then output root.
- Else if $\text{root} \rightarrow \text{low} > y$, then we recurse on $\text{root} \rightarrow \text{left}$.
- Else, $x > \text{root} \rightarrow \text{high}$.
 - If $\text{!empty}(\text{root} \rightarrow \text{left})$ and $x \leq \text{root} \rightarrow \text{left} \rightarrow \text{max}$, then there is an interval on the left containing x . We recurse on $\text{root} \rightarrow \text{left}$.
 - Else, we recurse on $\text{root} \rightarrow \text{right}$.

Complexity: $\mathcal{O}(\log n)$

Range Trees

Input: a static collection of k -dimensional points.

Recursive construction:

- An 1-range tree is a balanced binary search tree.
- A k -range tree, $k > 1$ is a balanced binary search tree on the first coordinate of each point.

Each node stores a $(k - 1)$ -dimensional range tree, for all remaining coordinates of all points in its rooted subtree.



Construction

- Find a point whose first coordinate is a median. – $\mathcal{O}(n)$.
- Construct k -range trees for left/right. – $2 \times C(n/2, k)$.
- Construct a $(k - 1)$ -range tree over the remaining coordinates and store it at the root. – $C(n, k - 1)$.

Complexity:

$$C(n, k) = C(n, k - 1) + 2 \times C(n/2, k) = \mathcal{O}(C(n, k - 1) \log n) = \mathcal{O}(n \log^k n)$$

Queries

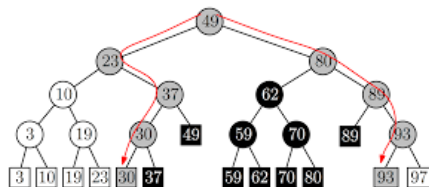
“Multi-dimensional range queries”

Input: a box = lower/upper bounds for all k coordinates

Output: enumerate all points in the box.

Variant: assign values to each point and output the max/min point, the sum of all values, etc.

A 1-dimensional range query with $[25, 90]$



Answer to a query (Sketch)

- 1) Consider the upper/lower bounds $[a_1, b_1]$ for the first coordinate.
- 2) In the binary search tree for the first coordinate, find the smallest/largest values x, y in the interval. Let $z = lca(x, y)$.

Remark: all searched points must be in the subtree rooted at z (otherwise, x or y could not be the smallest or largest value in $[a_1, b_1]$).

- 3) Consider the xz -path. For each edge uv on this path, if v is the left child of u , then all the right subtree of u contains points whose first coordinate lies between a_1 and b_1 .
- 4) Answer to queries on the remaining coordinates for the $(k - 1)$ -range trees on the paths between x, z and y, z .

Questions

