

- 1) Let v be an n -size vector. The q -quantiles of v are its kn/q order statistics, for $k = 1 \dots q$.
 - a. Show that we can compute all q -quantiles in $O(n \cdot \min(q, \log(n)))$.

We can sort the vector, that takes $O(n \cdot \log(n))$ time. Another possibility is to call q times (one for each quantile) the algorithm quickselect, and then the runtime is in $O(n \cdot q)$.

- b. Show that we can compute all q -quantiles in $O(n \cdot \log(q))$ time. Discuss about the optimality of this runtime.

We call quickselect for computing the kn/q order statistic, for $k = q/2$. It takes $O(n)$ time. In doing so, we also partitioned v in two subvectors v_L, v_R , that contain the elements smaller and larger than the kn/q order statistic, respectively. Now, half of the remaining q -quantiles is in v_L and the other half in v_R . Note that it implies that by applying recursively our algorithm to v_L, v_R , we can next compute two more q -quantiles, then four more q -quantiles, and so forth. In particular, the recursive depth until we computed all q -quantiles is in $O(\log(q))$. Since at each recursion level, all vectors considered are disjoint subvectors of v , the runtime at each recursion level is in total $O(n)$.

This is essentially optimal as, for $q=n$, computing all q -quantiles in order reduces to sorting.

Remark: the same approach works for computing any q order statistics (not necessarily the q -quantiles) in $O(n \cdot \log(q))$ time.

- 2) Consider a *sorted* n -size vector v and a *sorted* m -size vector u . The median of u and v , denoted in what follows by $\text{med}(u, v)$, is the median of the $(n+m)$ -size vector obtained from concatenating u and v together.
 - a. Show that $\text{med}(u, v)$ can be computed in $O(n+m)$ time.

We just concatenate u and v into one (not sorted) vector w , and then we apply to it quickselect. Note that this approach works even if u and/or v are not sorted...

- b. Show that if $m=n$, then $\text{med}(u, v)$ can be computed in $O(\log(n))$.

Let $\text{med}(u)$, $\text{med}(v)$ be the respective medians of u, v . We can compute both values in $O(1)$ because we assume u and v to be sorted. If $\text{med}(u) = \text{med}(v) = k$, then we also have $\text{med}(u, v) = k$. Otherwise, without loss of generality $\text{med}(u) < \text{med}(v)$. As we have $n/2$ elements from u , plus $> n/2$ elements from v , greater than $\text{med}(u)$, no element in the left half of u can be equal to $\text{med}(u, v)$. Similarly, no element in the right half of v can be equal to $\text{med}(u, v)$. So, let u_R, v_L be respectively the right half of u and the left half of v . We can apply recursively our algorithm to u_R, v_L since one can easily check that $\text{med}(u, v) = \text{med}(u_R, v_L)$. We stop the recursion as soon as there remains at most two elements in each vector, in which case we can solve the problem in $O(1)$. Since the size of the input vectors is halved at each recursion stage, the total runtime is in $O(\log(n))$.

- c. Show that $\text{med}(u, v)$ can be computed in $O(\log(\max(n, m)))$.

Without loss of generality, $m \neq n$. Let us consider the more general problem of computing the k th order statistic of u and v (for the median, $k = (n+m)/2$).

During a first stage of the algorithm, we essentially apply the previous approach of 2.b. More precisely, we first compute in $O(1)$ the medians $\text{med}(u)$, $\text{med}(v)$.

Case $\text{med}(u) = \text{med}(v)$. If $k = (n+m)/2$, then we are done (we output $\text{med}(u)$). If $k < (n+m)/2$, then we recurse on u_L, v_L with $k' = k$. If $k > (n+m)/2$, then we recurse on u_R, v_R with $k' = k - (n+m)/2$.

Case $\text{med}(u) < \text{med}(v)$. If $k = (n+m)/2$, then we recurse on u_R, v_L with $k' = k - m/2$. Note that we may not have $\text{med}(u, v) = \text{med}(u_R, v_L)$ in general, as $m \neq n$. If $k < (n+m)/2$, then we recurse on u, v_L with $k' = k$. If $k > (n+m)/2$, then we recurse on u_R, v with $k' = k - m/2$.

Case $\text{med}(v) < \text{med}(u)$. Similar to the previous case.

This above approach works until we can no longer make a recursive call. Indeed, when this happens, the smallest vector (say, u) stays with at most two elements, but the other vector may still have a non-constant (not even logarithmic) number of remaining elements, that prevents us from solving the problem in $O(1)$ during this last stage. So, instead, we do binary search for each of the (at most two) elements in u to determine where they should be inserted in the sorted vector v . In doing so, we know by how much we should shift the current position of the k th element in v in order to reach the k th order statistic of u, v . Note that we can shift either left or right (depending on where we should insert elements from u in v) and that we must shift by at most 2 positions in the vector.

Overall, the first stage where we can do recursive calls lasts $O(\log(m) + \log(n))$ steps. Indeed at each stage one of u, v has its size halved. The at most two binary searches of the last stage take $O(\log(\max\{n, m\}))$ time. The final runtime is also in $O(\log(\max\{n, m\}))$.

- d. Show that $\text{med}(u, v)$ can be computed in $O(\log(\min(n, m)))$.

Without loss of generality, $n < m$.

We slightly modify the previous approach from 2.c. First we propose a faster (but slightly less general) recursive stage than for 2.c. Let $\text{med}(u)$, $\text{med}(v)$ be the respective medians of u, v . If $\text{med}(u) = \text{med}(v) = k$, then we also have $\text{med}(u, v) = k$. Otherwise, without loss of generality $\text{med}(u) < \text{med}(v)$. Let u_R, v_L be respectively: vector u with its $n/2$ smallest element removed, and vector v with its $n/2$ biggest elements removed. Note that v_L is the lower half of v but that u_R is not the upper half of u . Furthermore, we can apply recursively our algorithm to u_R, v_L since one can easily check that $\text{med}(u, v) = \text{med}(u_R, v_L)$. *Indeed, we always remove the same number of elements to the left and to the right.* We stop the recursion as soon as there remains at most two elements in v . It takes $O(\log(n))$ recursive calls, and so, $O(\log(n))$ time.

Now, during the final stage recall that we need to shift by at most two positions in order to reach $\text{med}(u, v)$ from $\text{med}(u)$. Therefore, we do not need to know exactly where to insert the elements from v in u , but just to compare these at most two elements to the 5 middle elements of vector u . This can be done in

$O(1)$ time, that is much faster than binary search. Then, the total runtime is dominated by the first stage of the algorithm (where we still can do recursive calls), that is in $O(\log(n))$.

- 3) Consider two n -size vector v and w , such that all values in w are non-negative, and let $W = \sum_i w[i]$. A weighted median of v, w is a value x such that:

$$\begin{aligned} - \sum \{ w[j] \mid v[j] \leq x \} &\leq W/2 \\ - \sum \{ w[j] \mid v[j] > x \} &\leq W/2 \end{aligned}$$

- a. Propose a randomized algorithm for computing a weighted median, that achieves expected $O(n)$ time complexity.

We just use (randomized) quickselect. Each pivot (randomly selected) partitions the vector v in left/right subvectors v_L, v_R of smaller/larger elements, in $O(n)$. In doing so, we can compute the respective total weights W_L, W_R for both subvectors. If $\max\{W_L, W_R\} \leq W/2$, then our pivot is a weighted median. Else, if $W_L > W/2$, then we recurse on v_L (increasing the weight of its last element by $W - W_L$). Otherwise, we recurse on v_R (increasing the weight of its first element by $W - W_R$). The expected runtime of the algorithm follows the same recursive equation as for the classic version of quickselect (for unweighted median), and therefore it is also in $O(n)$.

- b. Propose a worst-case $O(n)$ -time algorithm for computing a weighted median.

Same as before but we use the deterministic version of quickselect ('median of medians').

- 4) Let M be an $n \times n$ matrix whose every column and row is sorted. Propose an $O(k \log(k))$ -time algorithm to compute the k th smallest element in the matrix.

We actually do something stronger: we compute all k smallest elements in the matrix sequentially. For that, at every step i , we maintain a set S_i of elements in the matrix (along with their coordinates in the matrix) such that the i th smallest element of the matrix is always in S_i ; in fact, it is always the smallest element in S_i . Initially, S_1 just contains the element $M[0,0]$, that is indeed the minimum element because we assume each row and column to be sorted. Then, to obtain S_{i+1} from S_i , we extract from S_i its smallest element v_i , that is by construction the i th smallest element from the matrix; let $v_i = M[a_i, b_i]$, then we replace v_i by $M[a_i+1, b_i]$ and $M[a_i, b_i+1]$ (if they exist). Note that $M[a_i+1, b_i]$ is the next element on the same (sorted) row as v_i , while $M[a_i, b_i+1]$ is the next element on the same (sorted) column as v_i .

If we naively stores S_i as a set, then the i th step of the algorithm would run in $O(|S_i|)$ time. By induction, $|S_i|$ is at most i , and therefore the total runtime would be linear in $1+2+\dots+i+\dots+k = O(k^2)$. Instead, all elements of S_i are stored in a self-balanced binary search tree. In doing so, each step runs in $O(\log(k))$ time.

Remark: here, it would be more natural to store S_i in a heap rather than in a binary search tree. Indeed, what we are doing here, but without telling it (because heaps were not presented in class at that point), is simulating a priority queue by using a self-balanced binary search tree.

- 5) In what follows, we are discussing about the optimal time for constructing a binary search tree (not necessarily balanced), being given n elements to be stored in this tree. We already know how to do this in $O(n \log n)$, e.g. by using AVL.
- Show that if we randomly insert these n elements in a binary search tree (using a naive implementation), then the expected runtime is also in $O(n \log n)$.

Let us assume that we choose the k th order statistic as the root (first element inserted in the tree). This happens with probability $1/n$. Then, $k-1$ elements are to be inserted to its left and the others to the right. In particular, all remaining $n-1$ elements are to be compared with the root when we insert them in the tree, and then we are reduced to constructing separately (at random) binary search trees for $k-1$ and $n-k$ elements, respectively. Therefore, the expected construction time satisfies $C(n) = n-1 + \frac{1}{n} \sum_k [C(k-1) + C(n-k)] = n-1 + \frac{2}{n} \sum_k C(k-1)$.

We conclude by using a classical trick in average analysis. More precisely, we have $(n+1)C(n+1) - nC(n) = n(n+1) - n(n-1) + 2C(n) = 2(n+C(n))$. Equivalently, $(n+1)C(n+1) = (n+2)C(n) + 2n$. Let $K(n) = C(n)/(n+1)$. We obtain $K(n+1) = K(n) + (2n)/((n+2)(n+1)) = K(n) + O(1/n)$. By induction, $K(n) = O(\sum_k 1/k) = O(\log n)$. As a result, $C(n) = O(n \log n)$.

- Show that if we are given n *sorted* elements, then we can construct a balanced binary search tree in $O(n)$ time.

We choose the median as the root of the tree, and then we apply the algorithm recursively to the elements that are smaller/larger than it to construct the left/right subtrees. Since the elements are sorted, finding their median and partitioning the remaining elements (depending on whether they are smaller or larger) can be done in $O(1)$. Therefore, the construction time satisfies $C(n) = O(1) + 2C(n/2) = O(n)$.

- Show that in general, the optimal runtime for constructing a binary search tree is in $O(n \log n)$. – Hint: Prove that you can use this tree in order to sort in less than $O(n \log n)$.

If we postorder the nodes in a binary search tree, then all values are sorted by increasing values. A postorder can be computed in $O(n)$. Therefore, the time to sort n elements is the time to construct a binary search tree + $O(n)$. It follows that constructing a binary search tree on n elements requires $\Omega(n \log n)$.

- 6) Consider an n -size balanced binary search tree T . The T -floor of an integer x is the greatest value stored in the tree that is smaller than or equal to x . Similarly, the T -ceil of an integer x is the least value stored in the tree that is greater than or equal to x . Show that we can compute the floor and the ceil of any value x in $O(\log n)$.

We search for x in T in $O(\log n)$. If x is stored in T , then both $\text{ceil}(x)$ and $\text{floor}(x)$ are equal to x . Otherwise, we compute $\text{ceil}(x)$ and $\text{floor}(x)$ during the search in T , as follows. Each time during the search we go on the left subtree of some node u , $\text{ceil}(x) = u.\text{value}$. Each time we go on the right subtree of some node u , $\text{floor}(x) = u.\text{value}$.

- 7) Consider an n -size balanced binary search tree T . Show that we can pre-process T in $O(n)$ such that the following types of queries can be answered in $O(\log(n))$:
- $q_1(x,y)$: what is the number of nodes u such that $x \leq u.value \leq y$?
 - $q_2(x,y)$: compute $\max \{ u.value \mid x \leq u.value \leq y \}$
 - $q_3(x,y)$: compute $\sum \{ u.value \mid x \leq u.value \leq y \}$

These are typical applications of range trees (recall that 1-range tree = balanced binary search tree). For each node u , we compute:

- $p_1(u)$: size of its subtree
- $p_2(u)$: maximum value stored in its subtree
- $p_3(u)$: sum of all values in its subtree.

It can be done in $O(n)$ by dynamic programming.

Now, to answer a query $q_3(x,y)$, we start computing nodes u and v such that: $u.value = \text{ceil}(x)$, $v.value = \text{floor}(y)$. As shown in the previous exercise, it can be done in $O(\log(n))$. Note that all values between x and y in T must be in fact between $u.value$ and $v.value$. Let $w = \text{lca}(u,v)$. Since T is balanced, it can be computed naively also in $O(\log(n))$. Furthermore, note that $u.value \leq w.value \leq v.value$.

- let $u_0=u, u_1, \dots, u_k=w$ be the uw -path in T (in the left subtree of w). We start summing all values $u_i.value$ between x and y . If furthermore u_i is the left child of u_{i+1} , then all values in the right subtree of u_{i+1} must be also between x and y ; in this situation, we add $p_3(u_{i+1} \rightarrow \text{right})$ to the total.

- let $v_0=v, v_1, \dots, v_k=w$ be the vw -path in T (in the right subtree of w). We start summing all values $v_i.value$ between x and y . If furthermore v_i is the right child of v_{i+1} , then all values in the left subtree of v_{i+1} must be also between x and y ; in this situation, we add $p_3(v_{i+1} \rightarrow \text{left})$ to the total.

The generalizations to the other two types of queries are straightforward.

- 8) Recall that in a 2-range tree (as seen in class), all values stored are bidimensional (2 coordinates $-x$ and $-y$). These values are stored in balanced binary search tree T_x , using for their keys only the $-x$ coordinates. Each node u in the tree stores a balanced binary search tree $T_y[u]$ where all values in its rooted subtree are stored, this time using for their keys the $-y$ coordinates. The construction time is in $O(n \cdot \log^2(n))$.

We consider the following modified construction of 2-range trees. Now, each node u of T_x stores all values in its subtree in an *array* $A_y[u]$ (no more a binary search tree), sorted by increasing $-y$ coordinates (breaking ties using $-x$ coordinates). If u is not a leaf, then let also u_L, u_R be its children nodes. For each element $A_y[u][i]$ we memorize the least indices i_L, i_R such that $A_y[u][i] \leq A_y[u_L][i_L]$ and $A_y[u][i] \leq A_y[u_R][i_R]$. Similarly, we memorise the largest indices j_L, j_R such that $A_y[u][i] \geq A_y[u_L][j_L]$ and $A_y[u][i] \geq A_y[u_R][j_R]$.

- Show that the construction time can be improved to $O(n \cdot \log(n))$.

We copy all values in two separate arrays, one B_x sorted by increasing $-x$ coordinates, and one B_y sorted by increasing $-y$ coordinates. It takes $O(n \cdot \log(n))$. Then, we do as follows: We choose the median (x^*, y^*) of B_x as root for T_x . Note that $A_y[x^*, y^*] = B_y$. Let $B_{L,x}, B_{R,x}$ be the values with smaller/larger $-x$ coordinate than the root (sorted by increasing $-x$ coordinates).

By scanning B_y once, we also construct in $O(n)$ time vectors $B_{L,y}$, $B_{R,y}$ of all values in $B_{L,x}$, $B_{R,x}$ sorted by $-y$ coordinates. Then, for each value $B_y[i]$, we compute the indices $i_L[i]$, $i_R[i]$ and $j_L[i]$, $j_R[i]$ in $B_{L,y}$, $B_{R,y}$ by using the following formulas:

- i) if $B_y[i] = B_{L,y}[i']$ for some i' then $i_L[i] = i'$.
- ii) Otherwise $i_L[i] = i_L[i+1]$.

We have similar formulas for all other three indices. As a result, all these values can be computed in $O(n)$ by dynamic programming.

We end up recursing on the disjoint pairs $B_{L,x}$, $B_{L,y}$ and $B_{R,x}$, $B_{R,y}$.

Recall that the initial sorting stage takes $O(n \cdot \log(n))$. The cost $C(n)$ of the recursive stage satisfies $C(n) = O(n) + 2 \cdot C(n/2)$. Hence, $C(n) = O(n \cdot \log(n))$.

- b. Let each value (x_i, y_i) be assigned some weight $w(x_i, y_i)$. Show that after a pre-processing in $O(n \cdot \log(n))$ time, we can answer in $O(\log(n))$ to the following type of queries $q(x_{\min}, x_{\max}, y_{\min}, y_{\max})$: output $\max \{ w(x_i) \mid x_{\min} \leq x_i \leq x_{\max} \text{ and } y_{\min} \leq y_i \leq y_{\max} \}$.

Pre-processing: at every node u of T_x , let $W_y[u] = w(A_y[u])$ (vector of weights of all points in the rooted subtree of u , sorted by increasing $-y$ coordinates), and let us construct a Cartesian tree $CT_y[u]$ for $W_y[u]$. The runtime is linear in $A_y[u].\text{size}$, and so the total runtime is at most the time needed to construct the 2-range tree, that is $O(n \cdot \log(n))$.

Answer to a query. Let r be the root of T_x . Using binary search on $A_y[r]$, we can compute in $O(\log(n))$ time indices i_{\min}, i_{\max} such that all values with their $-y$ coordinates between y_{\min}, y_{\max} are exactly those between i_{\min}, i_{\max} . Then, we apply the same approach as for exercise 7 (see also the discussion about range queries in class) in order to localize all values whose $-x$ coordinate is between x_{\min}, x_{\max} . The result is a set of $O(\log(n))$ rooted subtrees, and of $O(\log(n))$ additional values. We check directly amongst the $O(\log(n))$ additional values which of them also have their $-y$ coordinate between y_{\min}, y_{\max} and thus from now on we only focus on the $O(\log(n))$ rooted subtrees of T_x . Note that, as we traverse T_x starting from its root r , we can compute for each visited node u the indices $i_{\min}[u]$, $i_{\max}[u]$ such that all values in $A_y[u]$ with their $-y$ coordinates between y_{\min}, y_{\max} are exactly those between $i_{\min}[u]$, $i_{\max}[u]$. For that, it suffices to use the four pointers i_L, i_R, j_L, j_R stored for each element of $A_y[u]$. In particular, if u denotes the root of one of our $O(\log(n))$ rooted subtrees, then we can use the Cartesian tree $CT_y[u]$ in order to determine in $O(1)$ what the largest value stored in $W_y[u]$ between $i_{\min}[u]$, $i_{\max}[u]$ is.