

- 1) Let us consider the following variants of join/split operations.
 - a. Define a $\text{join}(T_1, T_2)$ operation, where all elements of T_1 are smaller than any element of T_2 . Note that we do not give any value i for the root.

We may assume T_1 to be nonempty. We compute the maximum element i from T_1 , and then we remove i from T_1 . Finally, we call $\text{join}(i, T_1, T_2)$. For a self-balanced binary search tree implementation, the runtime is in $O(\log(n))$.

- b. Define a $\text{split}(x, T)$ operation, with x possibly not in T .

We may assume x not in T (otherwise, this is the classical split operation on binary search trees). There are two cases:

Case $y > \min(T)$. Then, $y = \text{floor}(x, T)$ exists. Let T_1, T_2 be the output of $\text{split}(y, T)$. Finally, let T_1' be obtained from T_1 by inserting y . We output T_1', T_2 .

Case $y < \min(T)$. Then, $y = \text{ceil}(x, T)$ exists. Let T_1, T_2 be the output of $\text{split}(y, T)$. Finally, let T_2' be obtained from T_2 by inserting y . We output T_1, T_2' .

For a self-balanced binary search tree implementation, the runtime is in $O(\log(n))$.

- 2) Propose an efficient data structure which dynamically maintains a rooted tree subject to the following operations:
 - $\text{insert}(u, v)$: add a new leaf with father v
 - $\text{remove}(u)$: deletes leaf u (undefined if u is not a leaf)
 - $\text{lca}(u, v)$: returns the lca of nodes u, v .

If only insertions are allowed, then the following approach works (already discussed in a previous seminar): stores at each node u its level and pointers $p(u, j)$ to its 2^j -closest ancestors. In this situation insert and lca can be done in $O(\log(n))$. To remove a leaf u is in $O(1)$ because there is no x, j such that $p(x, j) = u$.

The following variant for computing the lca was discussed in class (another variant can be found in the correction of a previous seminar).

i) Set $x := u, y := v$. Without loss of generality, x and y are on the same level (if not, then we can use the shortcuts in order find a new pair u, v so that it is the case).

ii) Set $j :=$ largest exponent so that $p(x, j)$ and $p(y, j)$ are defined.

While $p(x, j) \neq p(y, j)$ do:

$x := p(x, j), y := p(y, j)$

$j :=$ largest exponent so that $p(x, j)$ and $p(y, j)$ are defined

done

From now on, j is defined such that $\text{lca}(x, y)$ is at most at distance 2^j from x, y .

iii) Consider the following procedure $\text{lca}(x, y, j)$:

If $j = 0$ then either $x = y$ or x, y have the same parent node

Else if $p(x, j-1) = p(y, j-1)$ then call $\text{lca}(x, y, j-1)$

Else call $\text{lca}(p(x, j-1), p(y, j-1), j-1)$

Since each step runs in $O(1)$ and j decreases at each step, the runtime is also in $O(\log(n))$.

- 3) Propose an efficient data structure which dynamically maintains a rooted tree subject to the following operations:

- $\text{insert}(u,v)$: add a new leaf with father v
- $\text{remove}(u)$: deletes leaf u (undefined if u is not a leaf)
- $\text{size}(u)$: returns the size of the subtree rooted at u .

We use an Euler tour tree (augmented Euler tour stored in a self-balanced binary search tree), where each node of the BST (edge or loop) further retains the size of its rooted subtree *in the BST* (not in the tree T represented). In doing so, insertion and removal of a leaf can be done in $O(\log(n))$. We can also compute $\text{size}(u)$ in $O(\log(n))$, as follows:

- if u is the root, then we output n
- else, let $e_1 = (v,u)$ and $e_2 = (u,v)$, with v the father of u . We may map both edges in some Hash-table to their respective positions in the BST. In doing so (going backward until the root), we may simulate searching for e_1 and e_2 . By searching for both edges in the BST, we can compute the respective numbers n_1, n_2 of elements (edges and loops) smaller than e_1, e_2 in the augmented Euler tour. In particular, there are $k = n_2 - n_1 - 1$ elements between e_1 and e_2 . These k elements represent an augmented Euler tour of the subtree rooted at u . If $\text{size}(u) = s$, then there are s loops and $s-1$ edges, where each edge is repeated twice. As a result, $k = 3s-2$.

- 4) Propose an efficient data structure which dynamically maintains a rooted tree *and* a preordering of its nodes subject to the following operations:
- $\text{insert}(u,v)$: add a new leaf with father v
 - $\text{remove}(u)$: deletes leaf u (undefined if u is not a leaf)
 - $\text{preorder}(u)$: returns the preorder of node u

This is the same approach as for the previous question (encoding by an Euler tour tree). The main difference is that now, each node x in the BST stores the number $p(x)$ of edges of the form $(u \rightarrow \text{father}(u))$ in its rooted subtree (i.e., it does not include the other edges and loops in the size of its rooted subtree). In doing so, we can compute $\text{preorder}(u)$ in $O(\log(n))$, as follows:

- if u is the root, then we output 0.
- else, let $e = (v,u)$, with v the father of u . By searching for edge e in the BST, we can compute the number k of edges $(y \rightarrow \text{father}(y))$ that appear before e in the tour. Then, $\text{preorder}(u) = k+1$.

- 5) Consider an n -size vector v . Show that after a pre-processing in $O(n \cdot \log(n))$, we can answer in $O(\log(n))$ to the following type of queries $q(i,j)$: “compute the number of *distinct* values e such that $v[k] = e$ for some k between i and j .”
- Hint: use 2-range trees.

We scan the vector v and, each time we read an element $v[i]=e$, we maintain in a Hash-table H the last position i where e was found (by convention, before we see an element for the first time, its associated value in H equals -1). Furthermore, upon reading $v[i]$, we also create a point $(H[v[i]], i)$. Note that we only create n points. We put all these points in a 2-range tree. It takes $O(n \cdot \log(n))$. Now, to answer a query $q(i,j)$, it suffices to compute the number of such points (x,y) such that: $x < i$ and $i \leq y \leq j$. Indeed, such a point witnesses a value between positions i and j whose latest appearance (if any) was before index i .

- a. (Discussed in class). By using Mo’s trick, solve the above query problem in $O(n \cdot \sqrt{n})$ pre-processing time and $O(\sqrt{n})$ query time.

- i) We partition the vector in $O(\sqrt{n})$ blocks of size $O(\sqrt{n})$. Let B_0, B_1, \dots, B_q denote the blocks.
- ii) Let $a[i]$ denote the number of distinct elements in B_i . To compute this number, it suffices to scan the i^{th} block. Each time we read a new value, we check whether this value was already inserted earlier in some auxiliary Hash table. The runtime is in $O(\sqrt{n})$ per block, and therefore it is in $O(n)$ in total.
- iii) Then, let $b[i, j]$ be a list that contains all distinct elements in B_i that *are not in any of* $B_{i+1} \cup B_{i+2} \cup \dots \cup B_{j-1} \cup B_j$. For $i=j$, we put all distinct elements of block B_i in $b[i, i]$ (these elements were computed at the previous step ii). Then, in order to construct $b[i, j+1]$ from $b[i, j]$, it suffices to check for each element of $b[i, j]$ whether it is contained in B_{j+1} . Since we put all elements of block B_{j+1} in an auxiliary Hash-table (cf. step ii), this takes $O(1)$ per element in $b[i, j]$, and so at most $O(\sqrt{n})$. Overall, since there are n entries to compute, and that each entry can be computed in $O(\sqrt{n})$, the runtime is in $O(n \cdot \sqrt{n})$.
- iv) Finally, let $c[i, j]$ be a Boolean value, equal to 1 if and only if $v[i]$ is in B_j . We apply the "partial sum trick" in order to compute $d[i, j] = \sum c[i, j'], j' = 0 \dots j$.

Now, to answer a query $q(i, j)$, we proceed as follows.

- * Let $B_p, B_{p+1}, \dots, B_{p+t}$ denote all blocks that are fully between i and j .
- * Now, consider all values $v[k]$, $k \geq i$, that appear *before* the first block B_p . In the same way, consider all values $v[k]$, $k \leq j$, that appear *after* the last block B_{p+t} . We put all these values in an $O(\sqrt{n})$ -size auxiliary u . By using a Hash-table, we remove from u any repetition of an element (all values are now unique).
- * We further remove from u any element that also appears in a block. For that, for each element $v[k]$ in u , it suffices to check whether $d[k, p+t] - d[k, p-1] > 0$.
- * We output $u.size() + \sum b[p+k, p+t].size()$, $k=0 \dots t$.

- 6) Consider an n -node tree T . For each node v , let $q(v)$ denotes the number of pairs of nodes (x, y) such that: v is an ancestor of x and y , $d(v, x) < d(v, y)$, $x.val > y.val$.
 - a. Show that in $O(n)$ we can edge-partition T in trees T_1, T_2 of respective orders between $n/3$ and $2n/3$.

We compute a centroid x of T and consider the components C_1, C_2, \dots, C_k of $T \setminus x$. If the largest component, say it is C_1 , has order $\geq n/3$, then we put $T_1 := C_1 + x$ and $T_2 := T \setminus C_1$. Otherwise, let i_1 be the smallest index such that there are $\geq 2n/3$ nodes in the union of C_1, C_2, \dots, C_{i_1} . Note that there are $< 2n/3$ nodes in the union of C_{i_1}, \dots, C_k because otherwise the order of C_{i_1} would be $\geq n/3$. We put $T_1 := \cup \{x, C_1, C_2, \dots, C_{i_1-1}\}$ and $T_2 := \cup \{x, C_{i_1}, \dots, C_k\}$.

- b. Show that in $O(n \cdot \log(n))$ we can *edge*-partition T in $O(\sqrt{n})$ subtrees of order $O(\sqrt{n})$ each.

We apply the centroid decomposition (see question a) until all gotten subtrees have order between $\sqrt{n}/3$ and \sqrt{n} .

Note that in general it is not possible to partition the *nodes*. For instance, if we are given a star, then one subtree should contain the root, and then any other subtree should be reduced to a leaf. Therefore, we would obtain $O(n)$ subtrees in our partition, not $O(\sqrt{n})$.

Remark: It is possible to compute an edge-partition as above in $O(n)$ by dynamic programming.

- c. Deduce from the above that after an $O(n*\sqrt{n}*\log(n))$ preprocessing, we can answer to any query $q(v)$ in $O(\sqrt{n})$.

0) We construct the edge-partition of question c. Let T_0, T_1, \dots, T_q be the subtrees of this edge-partition.

Then, every node is associated to exactly *one* subtree to which it belongs. More precisely, each node v different than the root is associated to the subtree which contains the edge $(v, \text{father}[v])$.

The set of nodes associated to one subtree is called a block. In what follows, let B_0, B_1, \dots, B_q denote the blocks. By construction, there are $O(\sqrt{n})$ blocks of size $O(\sqrt{n})$.

i) Let M_0 be the $n \times \sqrt{n}$ matrix where $M_0[x, i]$ denotes the number of nodes y in block i such that: $x.\text{level} < y.\text{level}$, and $x.\text{val} > y.\text{val}$.

In the same way, let M_1 be the $\sqrt{n} \times n$ matrix where $M_1[i, y]$ denotes the number of nodes x in block i such that: $x.\text{level} < y.\text{level}$, and $x.\text{val} > y.\text{val}$.

Both matrices can be constructed naively in $O(n^2)$. However, this can be improved as follows: If we represent all nodes in a block i as 2-dimensional points $(\text{value}, \text{level})$, and put these points in a 2-range tree (that costs us $O(\sqrt{n}*\log(n))$), then we can compute $M_0[x, i]$ and $M_1[i, y]$ in $O(\log(n))$. Overall, we can construct M_0, M_1 in $O(n*\sqrt{n}*\log(n))$.

ii) Let now M_2 be the $\sqrt{n} \times n$ matrix where $M_2[i, v]$ denotes the number of pairs (x, y) where x is in block i , y is in the subtree rooted at v , $x.\text{level} < y.\text{level}$, and $x.\text{val} > y.\text{val}$. For any i , by using M_1 we compute all entries $M_2[i, v]$ by dynamic programming on T (partial sum trick). So, we can construct M_2 in $O(n*\sqrt{n})$.

iii) Let M_3 be the $n \times \sqrt{n}$ matrix where $M_3[v, i]$ denotes the number of pairs (x, y) where x is *both* in the subtree rooted at v *and* in the same block as v , y is in block i , $x.\text{level} < y.\text{level}$, and $x.\text{val} > y.\text{val}$. For every i, j , by using M_0 we can compute all entries $M_3[v, i]$, v in block j , by dynamic programming on the j th subtree of the edge-partition. It takes $O(\sqrt{n})$ time for j fixed. Overall, we can construct M_3 in $O(n*\sqrt{n})$.

Now, to answer a query $q(v)$, we sum all values $M_2[i, v]$ for blocks i in the subtree rooted at v (each node may retain the list of all blocks in its subtree). In doing so, we compute the number of pairs (x, y) with x in a block of the subtree, and y anywhere in this subtree. However, in doing so we miss all pairs (x, y) such that: x is in the subtree, but the block containing x is *not fully* in the subtree rooted at v . (The same situation occurs for vectors, where we may have a block starting/ending before/after the interval we consider)

Let $B(v)$ be the block containing v . We claim that every block, except maybe $B(v)$, is either fully in the subtree rooted at v , or disjoint from it. Indeed, let us consider any block B that intersects both the subtree rooted at v and its complementary subtree. This block must be corresponding to the subtree containing edge $(v, \text{father}[v])$. As a result, we computed all desired pairs (x, y) except maybe those with x in $B(v)$. If $B(v)$ is *not* fully in the subtree rooted at v , then we also sum $M_3[v, B(v)]$ with all values $M_3[v, i]$.

The runtime is in $O(\sqrt{n})$ since there are at most $O(\sqrt{n})$ blocks.

- d. Remark: the above problem for trees is inspired by the problem of computing inversions in a sub-vector. However, it is simpler because:
- * for vectors, we were given two inputs for a query: i and j , and we wanted to count the number of inversions *only* between i and j
 - * for trees, we only have one input for a query: node v , and we want to count ``inversions'' in the subtree rooted at v .

Hence, for vectors, the real equivalent of our problem for trees would be as follows: *We are given an n -size vector v , that we want to pre-process in order to answer as fast as possible to the following type of queries $q(i)$: ``count the number of inversions in the sub-vector $v[i..n-1]$ ''.*

➔ This problem can be solved faster than by using Mo's trick. Indeed, in $O(n \log(n))$ time, we can compute for every index i the number $\text{inv}[i]$ of indices $j > i$ such that $v[i] > v[j]$. Then, we apply a classic ``partial sum trick'', namely: let $u[i] = \sum \text{inv}[k], k=0..i$. Now to answer to a query $q(i)$, it suffices to output $u[n-1] - u[i] + \text{inv}[i]$. The pre-processing time is in $O(n \log(n))$, while the query time is in $O(1)$.

It is not clear whether our problem for trees can also be solved faster than by using Mo's trick. *If we only consider pairs (x, y) such that x is an ancestor of y* , then indeed we can do better:

- * We perform a DFS traversal. In doing so, we also compute a preordering. For each node v , let $[p(v); q(v)]$ be the interval such that the descendants of v are exactly the nodes whose preorder is between $p(v)$ and $q(v)$. (In particular, $p(v)$ is the preorder of node v). All these operations can be done in $O(n)$.

- * Then, for each node u , we create a 2-dimensional point $(p(u), u.\text{val})$. We put all these points in a 2-range tree. This can be done in $O(n \log(n))$.

- * Now, given any node x , we may compute in $O(\log(n))$ the number $\text{inv}(x)$ of its descendants y such that $x.\text{val} > y.\text{val}$. For that, it suffices to compute the number of points (a, b) in our 2-range tree such that: $p(x) \leq a \leq q(x); b < x.\text{val}$, that is just a classical range query.

- * Finally, for any node v , let $u[v]$ be the sum of all values $\text{inv}[x]$, x a descendant of v . We can compute all values $u[v]$ by dynamic programming. Indeed, if v is a leaf, then $u[v] = \text{inv}[v] = 0$. Otherwise, let x_1, x_2, \dots, x_d be the children of node v . We have that $u[v] = \text{inv}[v] + \sum u[x_i], i = 1 \dots d$. **To answer to a query $q(v)$, it now suffices to output $u[v]$.**

A better equivalent of the inversion problem for trees could be as follows:

$q(u, v)$: compute the number of pairs (x, y) such that

- * x, y are nodes of the unique path between u and v

- * $d(u,x) < d(u,y)$, i.e. u is closer to x and v is closer to y
- * $x.val > y.val$.

We can use Mo's trick in order to solve this above problem. But this is a bit more complicated...