

ER - DIAGRAM

LECTURE 1: Databases

DBMS (Database Management System)

- System designed to define and manipulate data.
 - Storage.
 - Retrieval.
 - Updates.

DBMS (Database Management System)

- Avoid redundancy, inconsistency.
- Concurrent data access.
- Provides security and recovery.
- Declarative language to manipulate, query, define and control data.
- DDL, DML, DCL.
- Data dictionary: database providing info about database structure.

DBMS (Database Management System)




- Text database, example **CSV** format.
- Implemented in 1970 (IBM).
- File = table with a single record on each line.
- Read, store and send.
- Simple structure.
- Inefficient: slow, duplicated values, difficult to update etc.

DBMS (Database Management System)




- Tree structure, examples: file system, Windows Registry
- IBM Information Management System (**IMS**)
- XML, XAML
- Used in mainframe era.
- Rigid structure.
- Only *One-to-many* relationship.
- Traversing very easy, moving a node difficult

DBMS (Database Management System)



- Hybrid relation + objects =>> tables of objects.
- **Realm** database for Android/iOS: classes used as schema definition, alternative for SQLite.
 - Lightweight, doesn't need an ORM
 - MVCC architecture (multiversion concurrency control)
- MongoDB Realm sync data [1].

DBMS (Database Management System)



- Transaction oriented systems (example financial transactions).
- **ACID**: Atomicity, consistency, isolation, durability.
- Suitable for structured data.

DBMS (Database Management System)



- RDBMS hard to scale (scale vertically, not horizontally).
- RDB Restrictive schemas =>> flexible structure.
- The state of the database can change.
- **!!!** availability, scalability, performance
- Sharding: distribute data on different servers

DBMS (Database Management System)



- Cloud and bigdata.
- **BASE** (Basically Available, Soft state, Eventually consistent)

Database appears to work most of the time

Replicas are not consistent all the time

At some later point database will be consistent.

DBMS (Database Management System)



- Cloud and bigdata.
- **BASE** (Basically Available, Soft state, Eventually consistent)
- Types:
 - key-value: Redis
 - Document: Mongo, Firestore
 - Column: Apache Cassandra, HBase, MariaDB
 - Graph: Neo4j

Sql or NoSQL

Relational

- Vertical scalability
- Consistency model: ACID
- pre-defined schema
- SQL language
- Normalized data

NoSql

- Horizontal scalability
- Consistency model: BASE
- Flexible schema
- No standard
- Collections, redundancy

DBMS (Database Management System)




- Integration of Relational and NoSQL databases.
- Integration of in-memory DB and on-disk DB
- Altibase, Orient DB

Course roadmap

- Database design (ERD, Relational model, normalization)
- Transactional systems, consistency models, concurrency control
- NoSql and BidData topics
- SQL (LMD, LDD, LCD)
- ... & other topics ...
- Please answer [www.menti.com](https://www.menti.com/join/80457859) 8045 7859 Q1, Q2, Q3, Q4, Q5

Course roadmap



ER - Diagram

Entity Relationship model

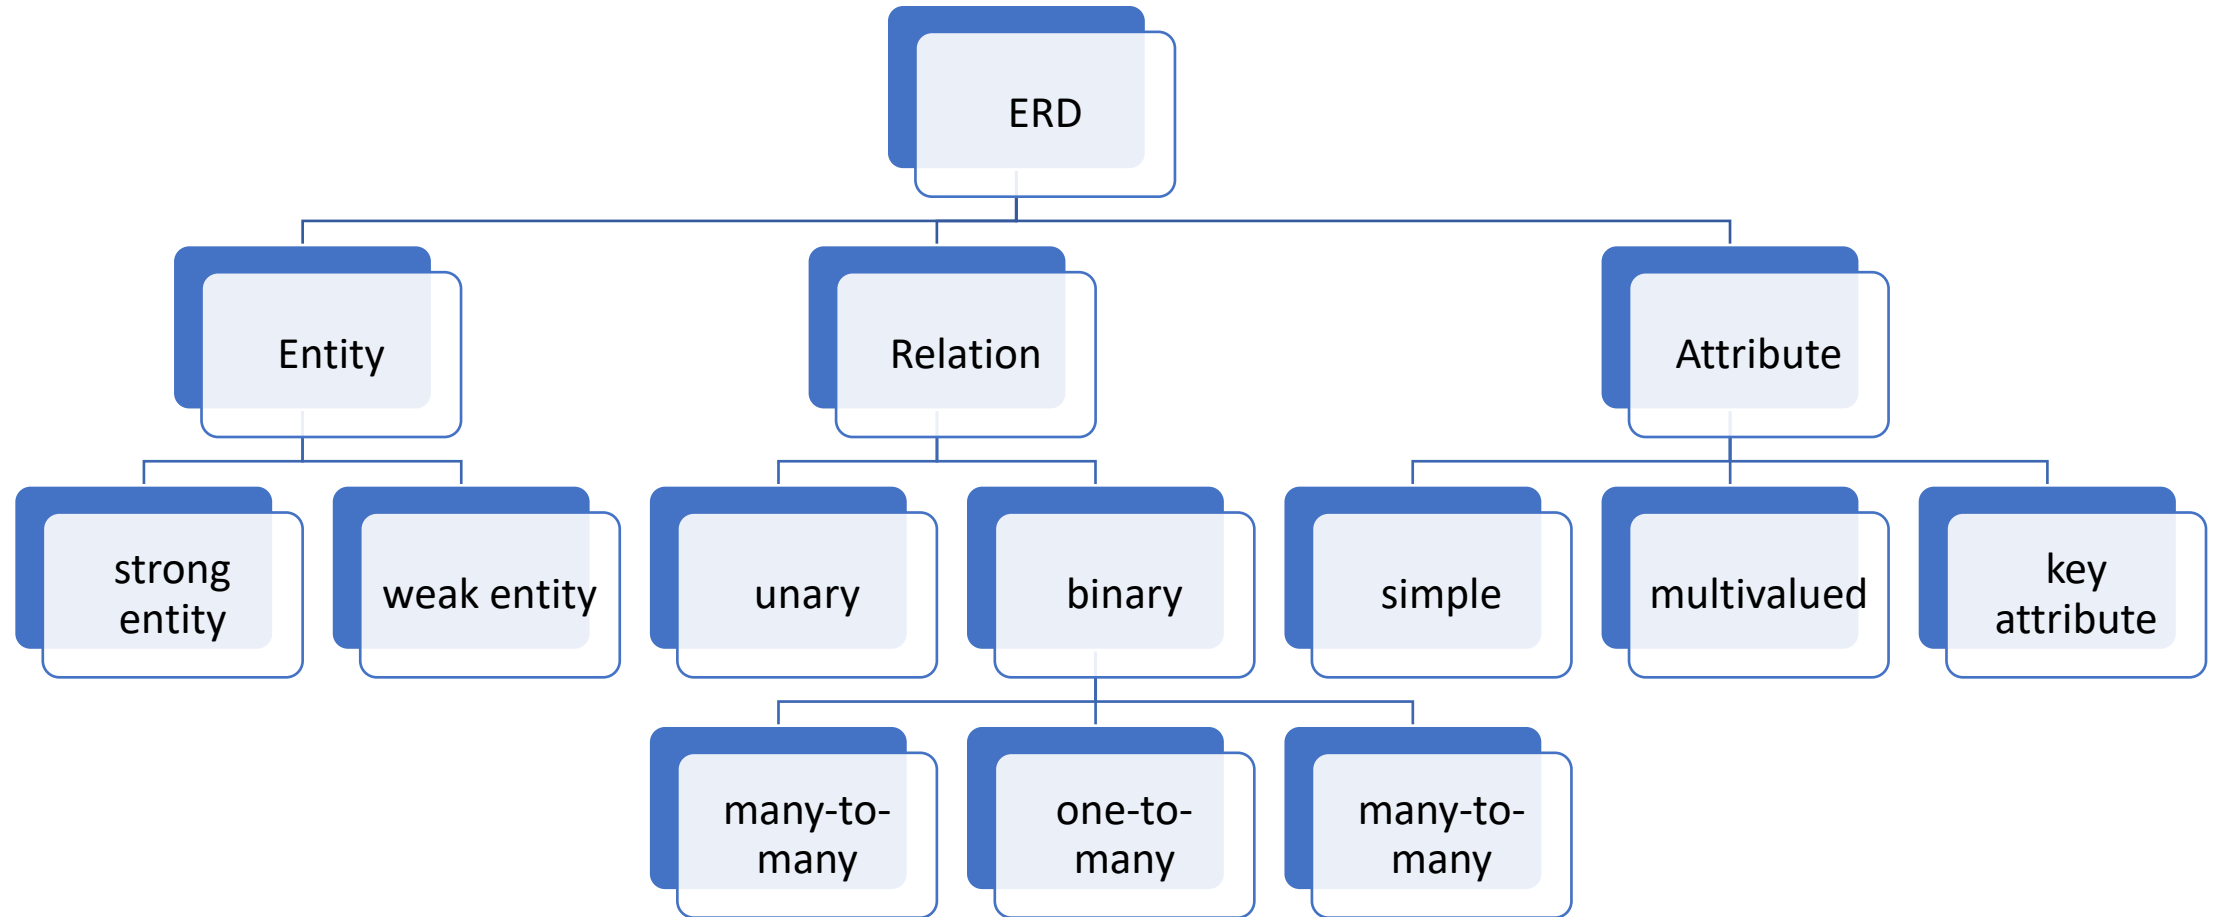
ER - Diagram

- Visual representation of the ER conceptual data model.
- High-level design.
- Not linked to the implementation or hardware.
- Peter Chen proposed ERDs in 1976.

ER - Diagram

- User story/requirement analysis → **ER** → relational database schema.
- Easy to translate into relational tables.
- Describes the logical structure of the (relational) database.
- Suitable for structured systems.

ERD - components



ER - Diagram



person, place, activity, event, concept, real world object etc.
usually a noun



ER - Diagram



ENTITY

person, place, activity, event, concept, real world object etc.
usually a noun



RELATION

links entities (unary, binary, ternary).
usually a verb



ATTRIBUTE

ER - Diagram



ENTITY

person, place, activity, event, concept, real world object etc.
usually a noun



RELATION

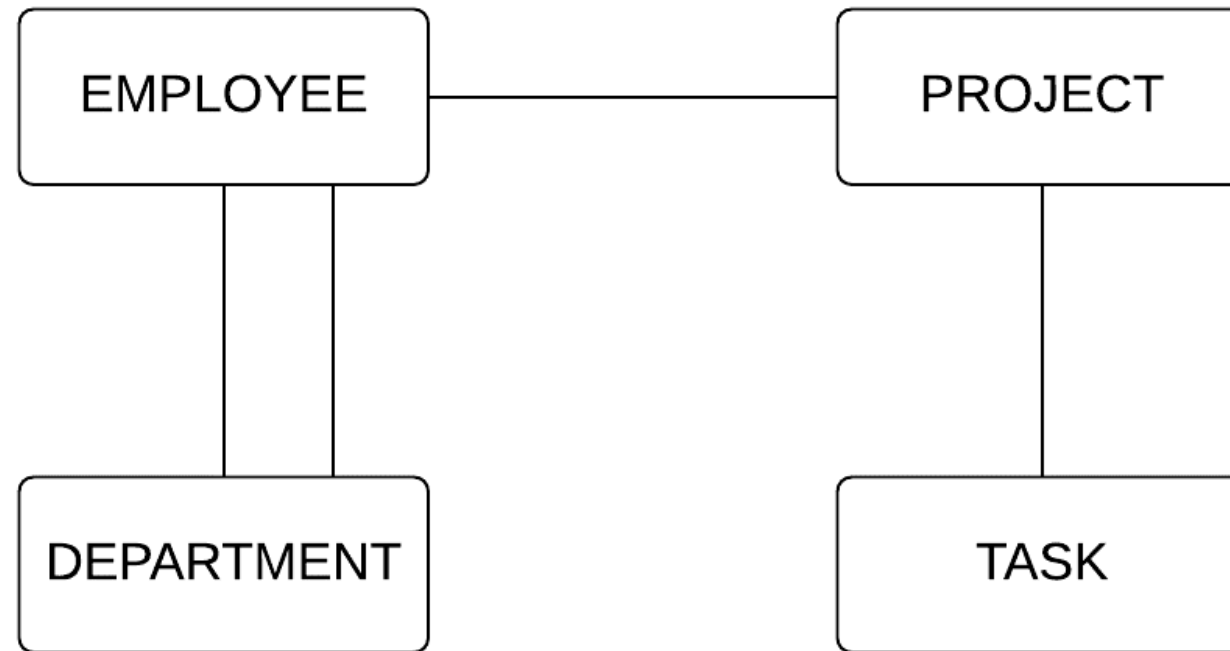
links entities (unary, binary, ternary).
usually a verb



ATTRIBUTE

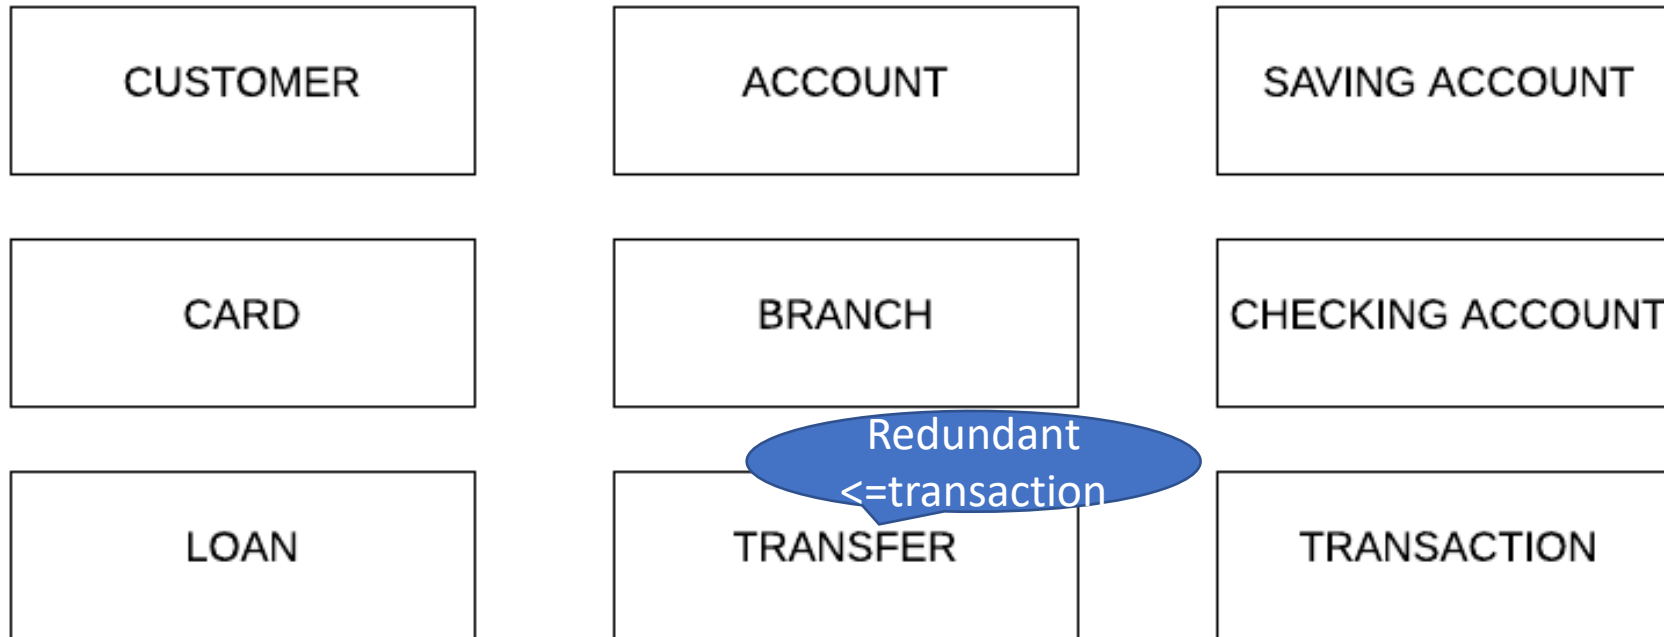
describe entities or relations


Entities

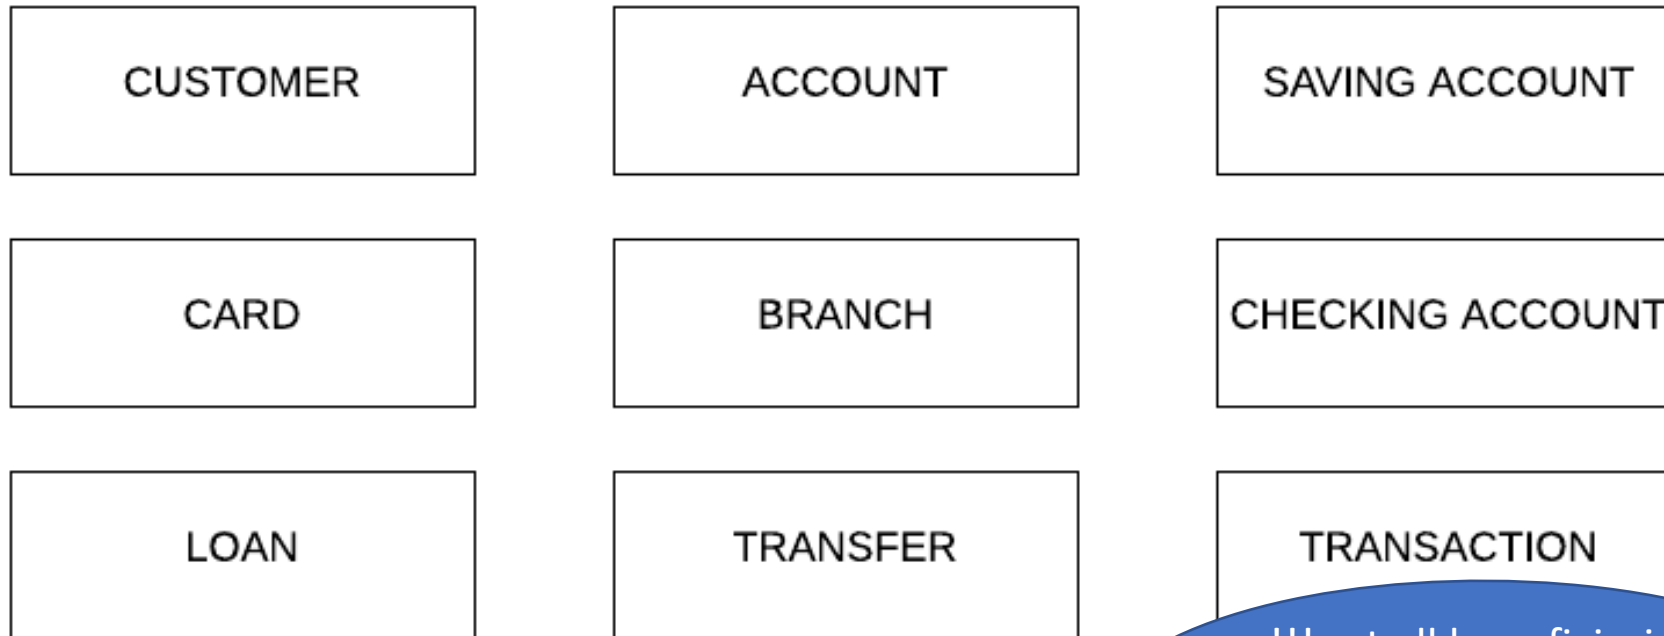


Banking (1)Entities

- A customer opens a saving account or a checking account, at a bank branch. He may also access loans. For each checking account he has a card. Periodically he may withdraw money from his account or partially pay his loans. He may also transfer money from one account to another.
- Please answer www.menti.com 8045 7859 Q6, Q7







!!!not all beneficiaries
(missing from story) are
customers of the same
bank

Entities

- Unique names, uppercase characters
- Graphical representation: rectangles
- Relational database: entity ➔ table (line & columns)
- Primary key: attribute or group of attributes that uniquely identifies an entity instance

Primary key

- **Unique** identifier
 - Must be known at any moment (**not null**)
 - Simple
 - No ambiguities
 - Immutable
-
- Composed keys may be replaced with an *artificial key*.
 - In many RDBMS we may use autoincremented values.

Primary key UUID/GUID

- **universally unique identifier** 128-bit
 - Probability of collision (that a UUID is duplicated) is negligible.
 - No need to change when merging to databases.
 - Known before the insertion of a new row, without querying the database.
-
- Please answer [www.menti.com](https://www.menti.com/join/80457859) 8045 7859 Q8

Primary key UUID/GUID

- **universally unique identifier** 128-bit
- Not the best solution for clusters (use sequential UUIDs).
- Types:
 - Type 1 : 4 bytes + 2 bytes + 2 bytes + 2 bytes + 6 bytes = time + node
 - Type 4 : 122 bits randomly generated, 6 bits reserved for version and variant.
- Bit for type
 - type 1 2ad1db02-2ff0-11eb-**a**dc1-0242ac120002
 - type 4 a7bc2d72-7153-44a1-**8**3df-d03dd298cf53

Candidate key

- *Relation* in relational model \leftrightarrow *relationship* in ERD
- Relation \leftrightarrow set of tuples \leftrightarrow tables \leftrightarrow attributes of relation are columns
- Minimal *superkey* or *o relation* (minimal set of attributes) such that:
 - 1) There are no two distinct tuples sharing the same values for the *superkey* (unique)
 - 2) No proper subsets of the superkey has property (1)
- Attributes of the superkey are called prime attributes.
- Attributes that does not occur in any superkey are non-prime attributes.
- If there are no null values, since a relation is a *set* of tuples, each relation has at least one candidate key, i.e the set of all its attributes.

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

There are no candidate key with one attribute

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Unique keys: ~~(BUYER_ID, PRICE)~~

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Unique keys: ~~(BUYER_ID, HALL_NO)~~

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Unique keys: (BUYER_ID, DATE) (BUYER_ID, VIP) (PRICE, DATE) (HALL_NO, DATE) (HALL_NO, TYPE)
(BUYER_ID, PRICE, HALL_NO) (BUYER_ID, HALL_NO, DATE)
etc -- all sets with 3, 4 or 5 attributes

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Candidate keys: (BUYER_ID, DATE) (BUYER_ID, VIP) (PRICE, DATE) (HALL_NO, DATE) (HALL_NO, TYPE)
(BUYER_ID, PRICE, HALL_NO) (BUYER_ID, PRICE, TYPE) ,
~~(PRICE, DATE, DATE, TYPE)~~

Candidate keys

BUYER_ID	PRICE	HALL_NO	DATE	TYPE
1	150	Coliseum	08/03/22	VIP
1	150	Lyttelton	14/04/22	A
2	140	Olivier	01/05/22	A
2	90	Coliseum	04/06/22	B
2	220	Lyttelton	08/03/22	VIP
3	140	Olivier	14/04/22	B
3	220	Olivier	20/03/22	VIP

Candidate keys: (BUYER_ID, DATE) (BUYER_ID, VIP) (PRICE, DATE) (HALL_NO, DATE) (HALL_NO, TYPE)
(BUYER_ID, PRICE, HALL_NO) (BUYER_ID, PRICE, TYPE) ,

Participant
participant_id last_name first_name

Product
product_id name code reserve_price

Info
photo description

Bid
price timestamp participant_id product_id currency

candidate
keys?

Category
category_id name

<<enum>> Currency
USD EUR GBP ...

Please answer www.menti.com 44 32 42 0

Q8



Airline (1)Entities

- The airline has one or more airplanes. An airplane has a model number, and capacity. Each flight is carried out by airplanes. An airplane is uniquely identified by its Registration_no and a flight is identified by its Flight_no. A passenger can book a ticket for a flight.

Please answer www.menti.com 8045 7859

Q10, Q11



Entities



- Weak entity is an entity that depends on another entity.
- The primary key of a weak entity contains the primary key of the strong entity that it depends on + description/partial key.

Please answer www.menti.com 8045 7859 Q12, 13

Entities



- A sub-entity has the same key as the *super*-entity and all its attributes and relationships.

Please answer www.menti.com 8045 7859

Q14

Airline (1) Relationships

- The airline has one or more airplanes. An airplane has a model number, and capacity. Each flight is carried out by airplanes. An airplane is uniquely identified by its Registration_No and a flight is identified by its Flight_No. A passenger can book a ticket for a flight.

Airline (1) Relationships

- The airline has one or more airplanes. An airplane has a model number, and capacity. Each flight is carried out by airplanes. An airplane is uniquely identified by its Registration_No and a flight is identified by its Flight_No. A passenger can book a ticket for a flight. A flight may have **one or more stops**.

Airline (1) Relationships

- The airline has one or more airplanes. An airplane has a model number, and capacity. Each flight is carried out by airplanes. An airplane is uniquely identified by its Registration_No and a flight is identified by its Flight_No. A passenger can book a ticket for a flight. A flight may have one or more stops. The passenger will pay for **extra baggage**.

- [\[1\] https://docs.mongodb.com/realm/node/sync-data/](https://docs.mongodb.com/realm/node/sync-data/)
- [\[2\] https://tomharrisonjr.com/uuid-or-guid-as-primary-keys-be-careful-7b2aa3dcb439](https://tomharrisonjr.com/uuid-or-guid-as-primary-keys-be-careful-7b2aa3dcb439)
- [\[3\] https://database.guide/newid-vs-newsequentialid-in-sql-server-whats-the-difference/](https://database.guide/newid-vs-newsequentialid-in-sql-server-whats-the-difference/)
- [\[4\] https://docs.mongodb.com/realm/sdk/android/advanced-guides/threading/#std-label-android-mvcc](https://docs.mongodb.com/realm/sdk/android/advanced-guides/threading/#std-label-android-mvcc)